



INNOVATE2018

ONLINE CONFERENCE

DEVELOPER EDITION

Introduction to Deep Learning Theory, Use Cases, and Tools

Julio Faerman @faermanj

Julien Simon @julsimon

@awscloud



#AWSInnovate

Myth: AI is dark magic

aka “you’re not smart enough”



What to expect

1. An introduction to Deep Learning
2. Common network architectures and use cases
3. Resources for diving deeper

Artificial Intelligence: design software applications which exhibit human-like behavior, e.g. speech, natural language processing, reasoning or intuition

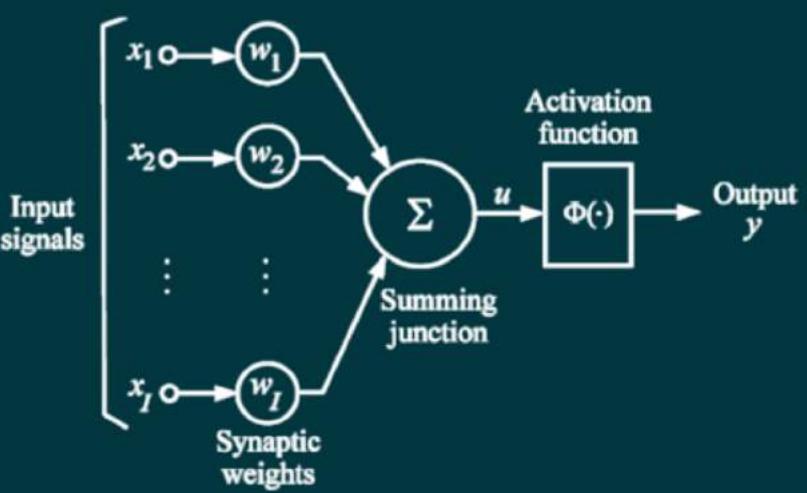
Machine Learning: teach machines to learn without being explicitly programmed

Deep Learning: using neural networks, teach machines to learn from complex data where features cannot be explicitly expressed



Fact: AI is math, code and chips

A bit of science, a lot of engineering



```
data = mx.symbol.Variable('data')
conv1 = mx.sym.Convolution(data=data, kernel=(5,5), num_filter=20)
relu1 = mx.sym.Activation(data=conv1, act_type="relu")
pool1 = mx.sym.Pooling(data=relu1, pool_type="max", kernel=(2,2), stride=(2,2))
conv2 = mx.sym.Convolution(data=pool1, kernel=(5,5), num_filter=50)
relu2 = mx.sym.Activation(data=conv2, act_type="relu")
pool2 = mx.sym.Pooling(data=relu2, pool_type="max", kernel=(2,2), stride=(2,2))
flatten = mx.sym.Flatten(data=pool2)
fc1 = mx.symbol.FullyConnected(data=flatten, num_hidden=500)
relu3 = mx.sym.Activation(data=fc1, act_type="relu")
fc2 = mx.sym.FullyConnected(data=relu3, num_hidden=10)
lenet = mx.sym.SoftmaxOutput(data=fc2, name='softmax')
```



ML / DL hardware ?

Google TPU, Intel Nervana, Xilinx Everest

Quantization: using integer or binary weights and activations

- Reduces power consumption
- Simplifies the logic needed to implement the model
- Reduces memory usage

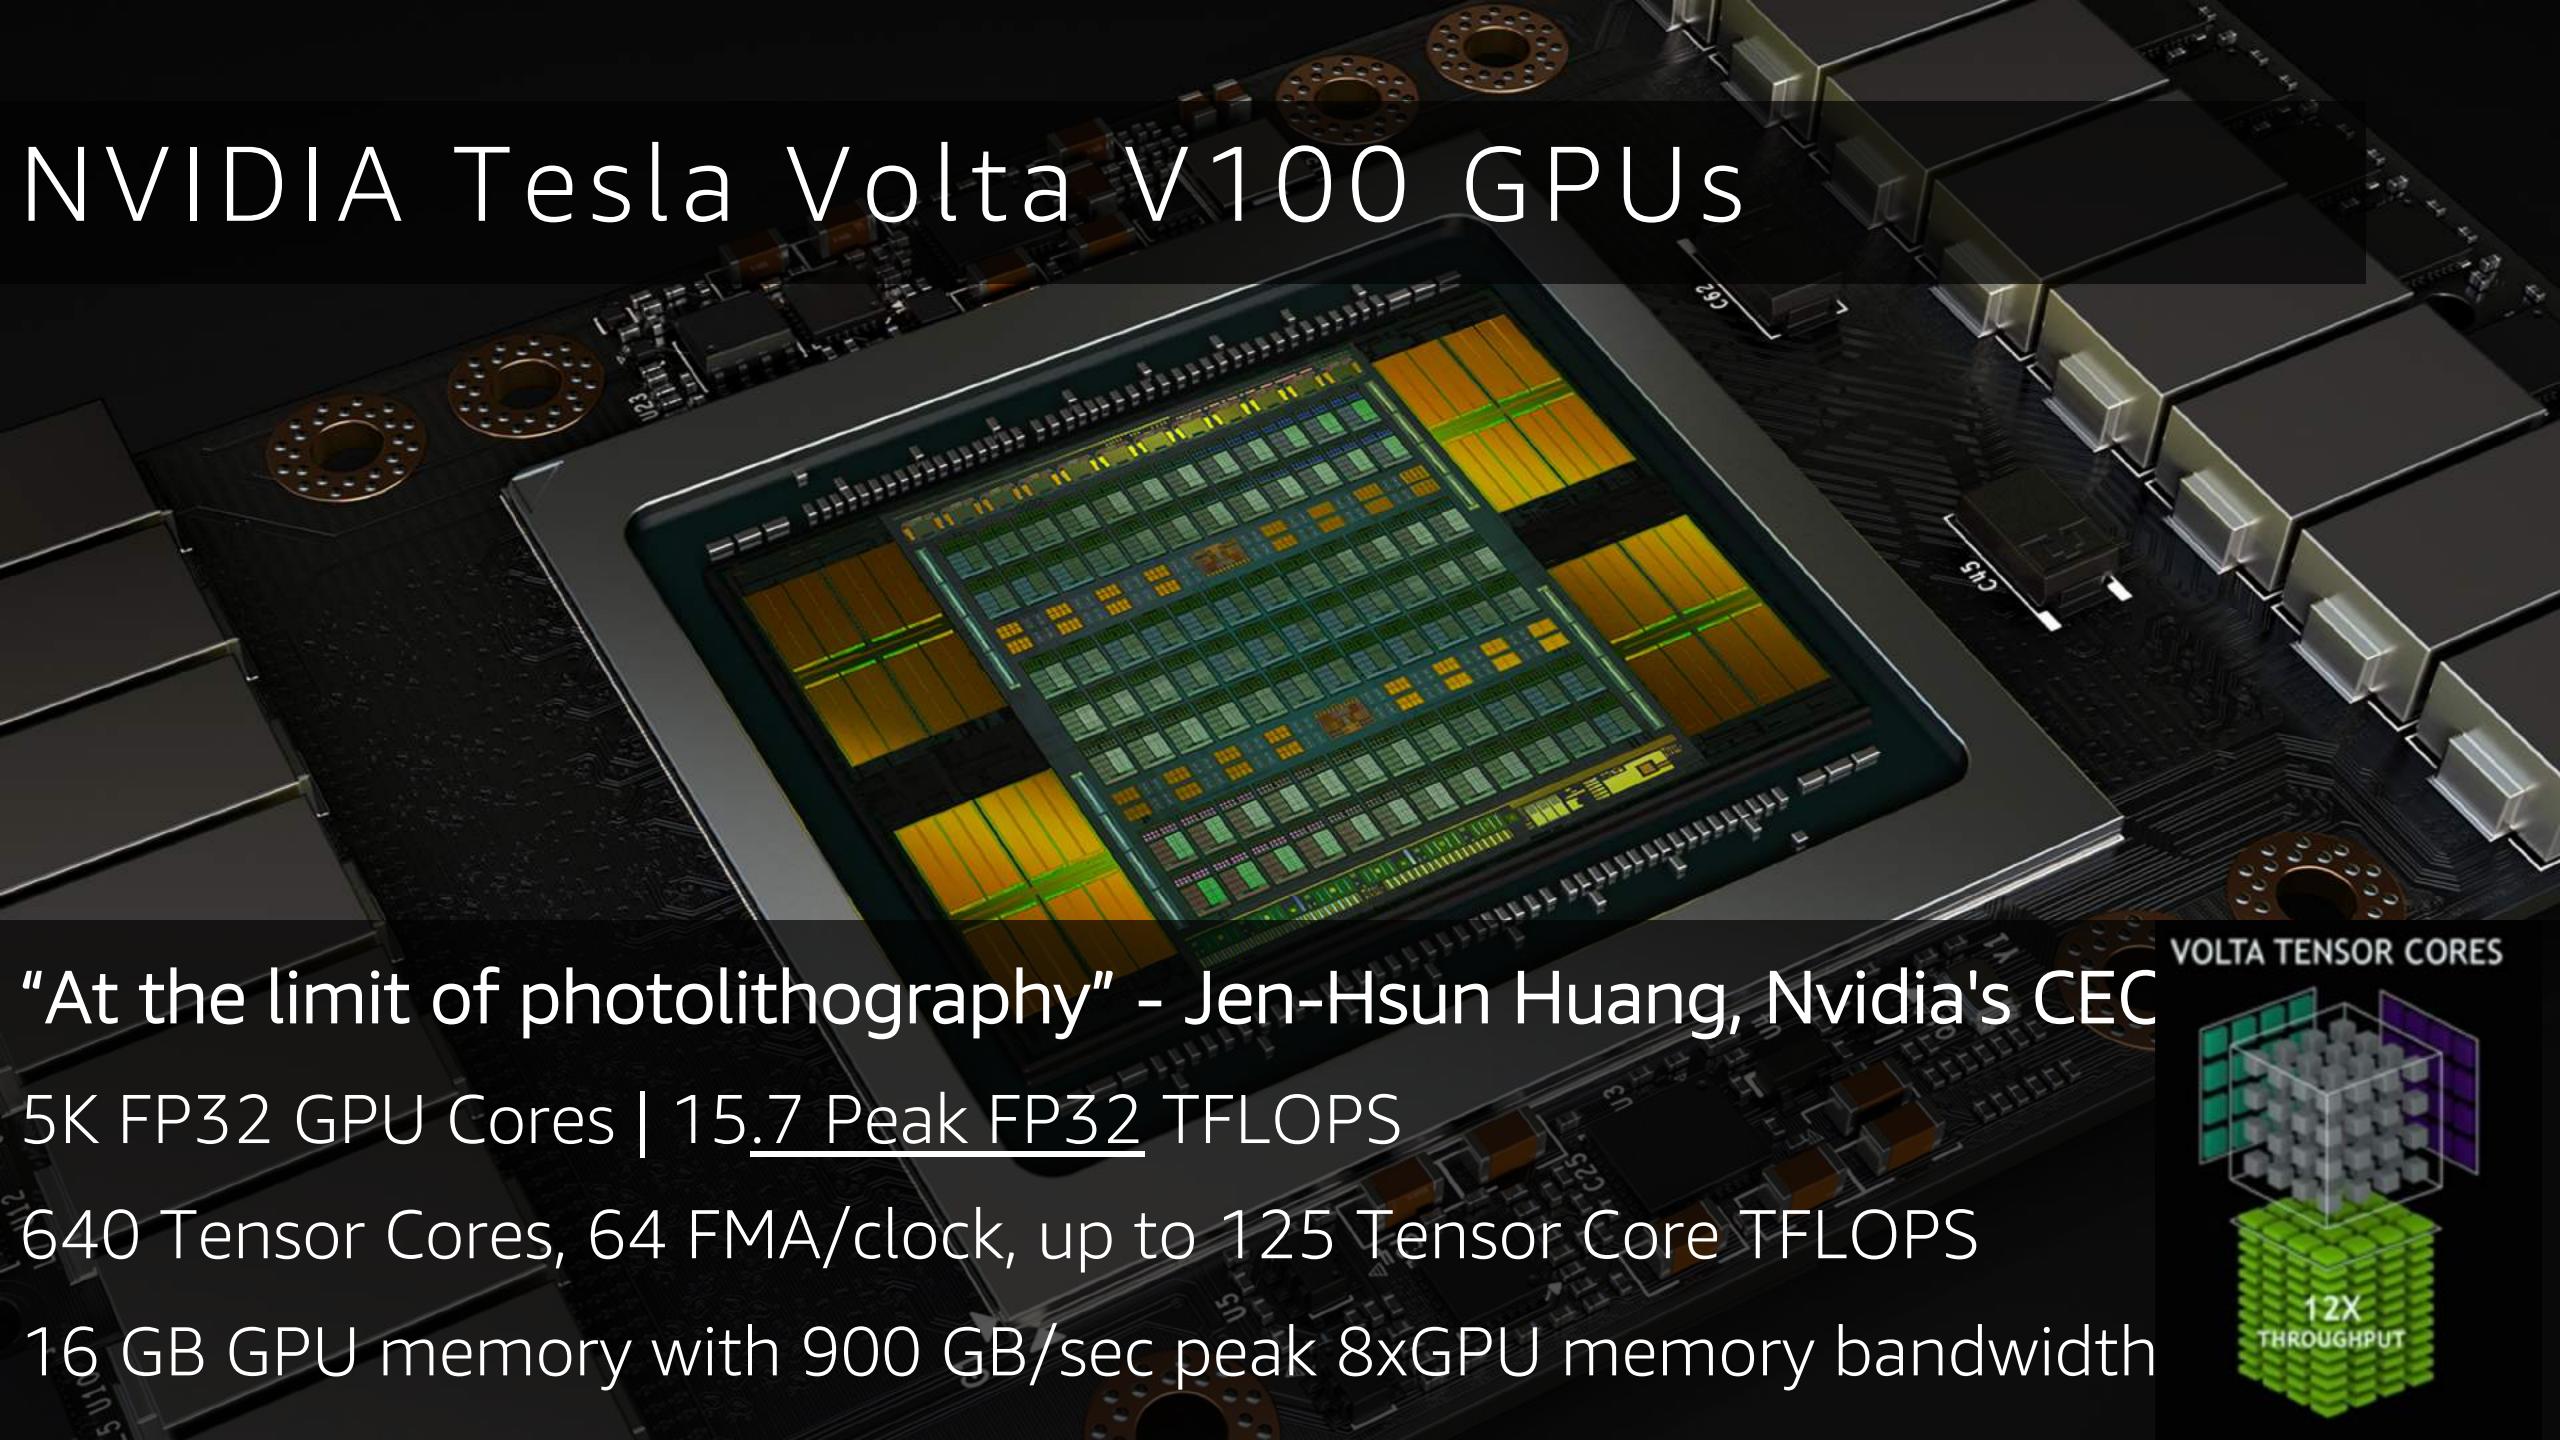
Pruning: removing useless connections

- Increases computation speed
- Reduces memory usage

Compression: encoding weights

- Reduces model size

NVIDIA Tesla Volta V100 GPUs



"At the limit of photolithography" - Jen-Hsun Huang, Nvidia's CEC

5K FP32 GPU Cores | 15.7 Peak FP32 TFLOPS

640 Tensor Cores, 64 FMA/clock, up to 125 Tensor Core TFLOPS

16 GB GPU memory with 900 GB/sec peak 8xGPU memory bandwidth



Netflix Technology Blog

[Follow](#)

Learn more about how Netflix designs, builds, and operates our systems and engineering organizations

Feb 10, 2014 · 11 min read

Distributed Neural Networks with GPUs in the AWS Cloud

by [Alex Chen](#), [Justin Basilico](#), and [Xavier Amatriain](#)

As we have described [previously](#) on this blog, at Netflix we are constantly innovating by looking for better ways to find the best movies and TV shows for our members. When a new algorithmic technique such as Deep Learning shows promising results in other domains (e.g. [Image Recognition](#), [Neuro-imaging](#), [Language Models](#), and [Speech Recognition](#)), it should not come as a surprise that we would try to figure out how to apply such techniques to improve our product. In this post, we will focus on what we have learned



ML @ AWS

Put machine learning in the hands of every developer and data scientist





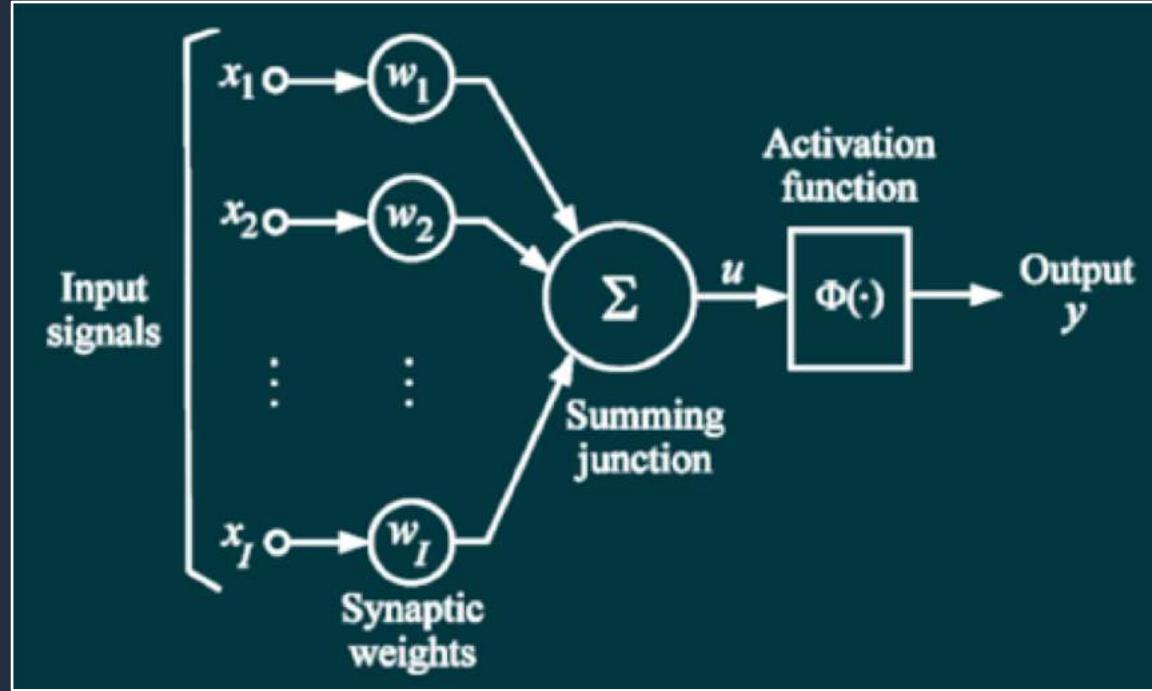
INNOVATE2018

ONLINE CONFERENCE

DEVELOPER EDITION

An introduction to
Deep Learning

The neuron



$$\sum_{i=1}^l x_i * w_i = u$$

"Multiply and Accumulate"

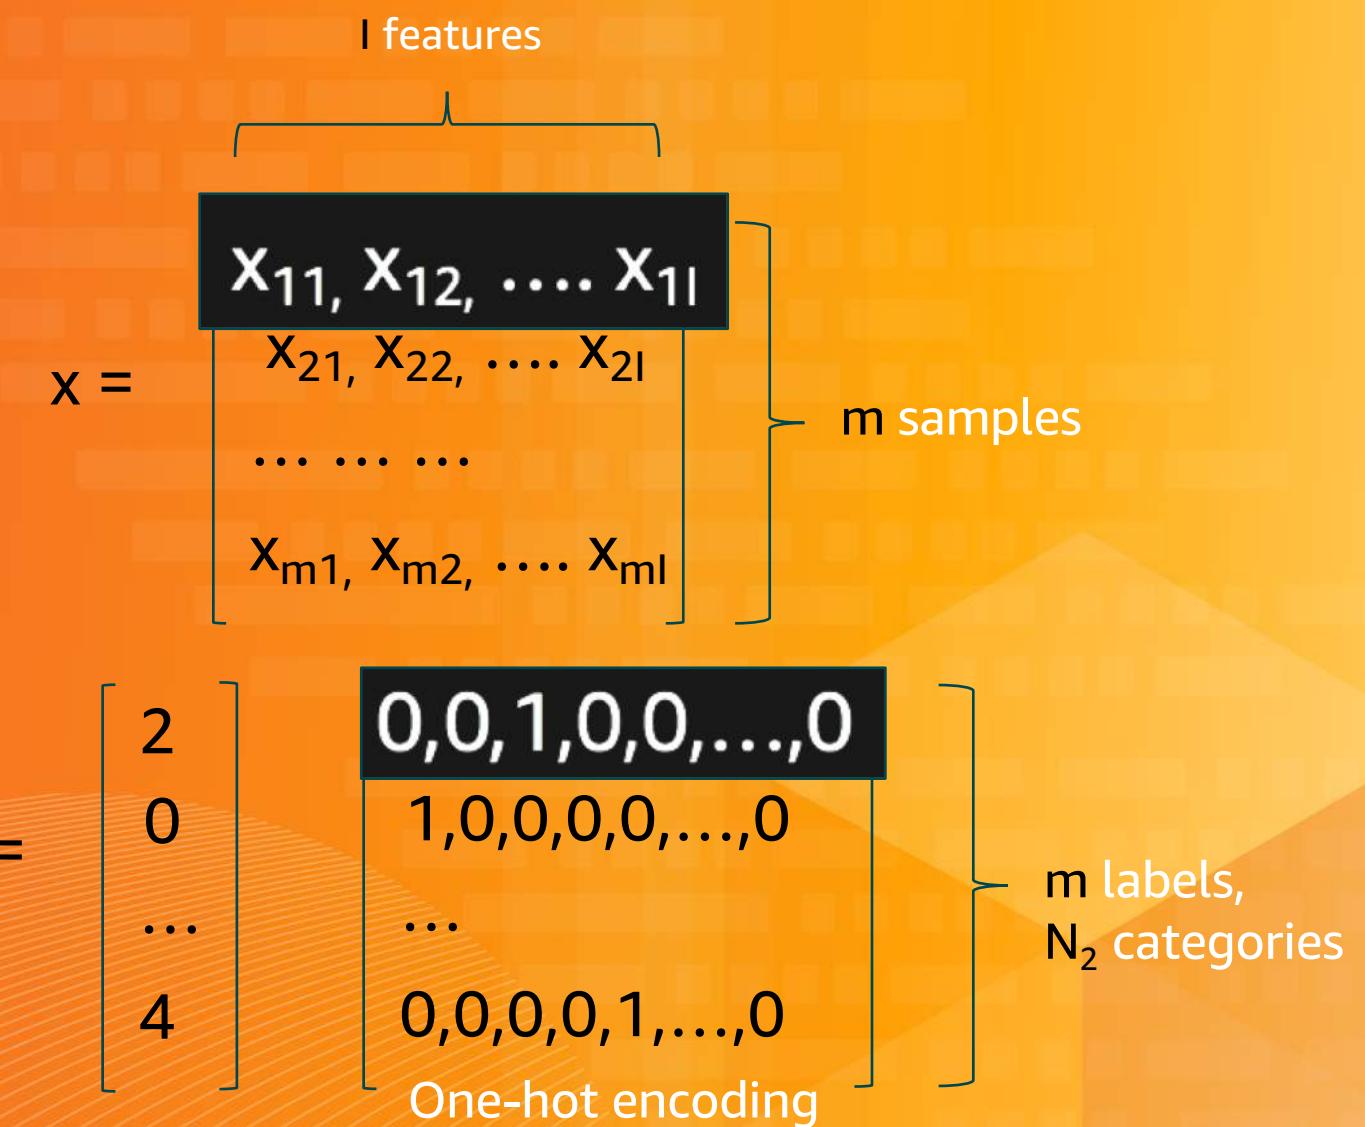
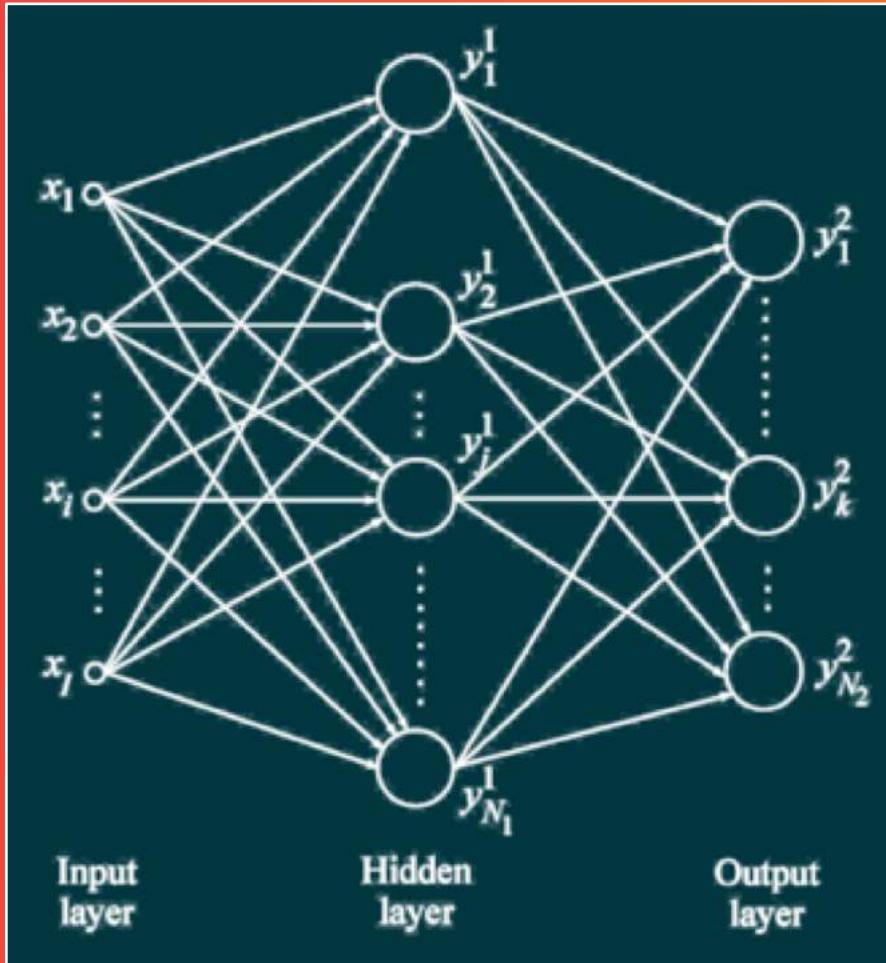
Activation functions

Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a. Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Softsign [7][8]		$f(x) = \frac{x}{1 + x }$
Rectified linear unit (ReLU) [9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$

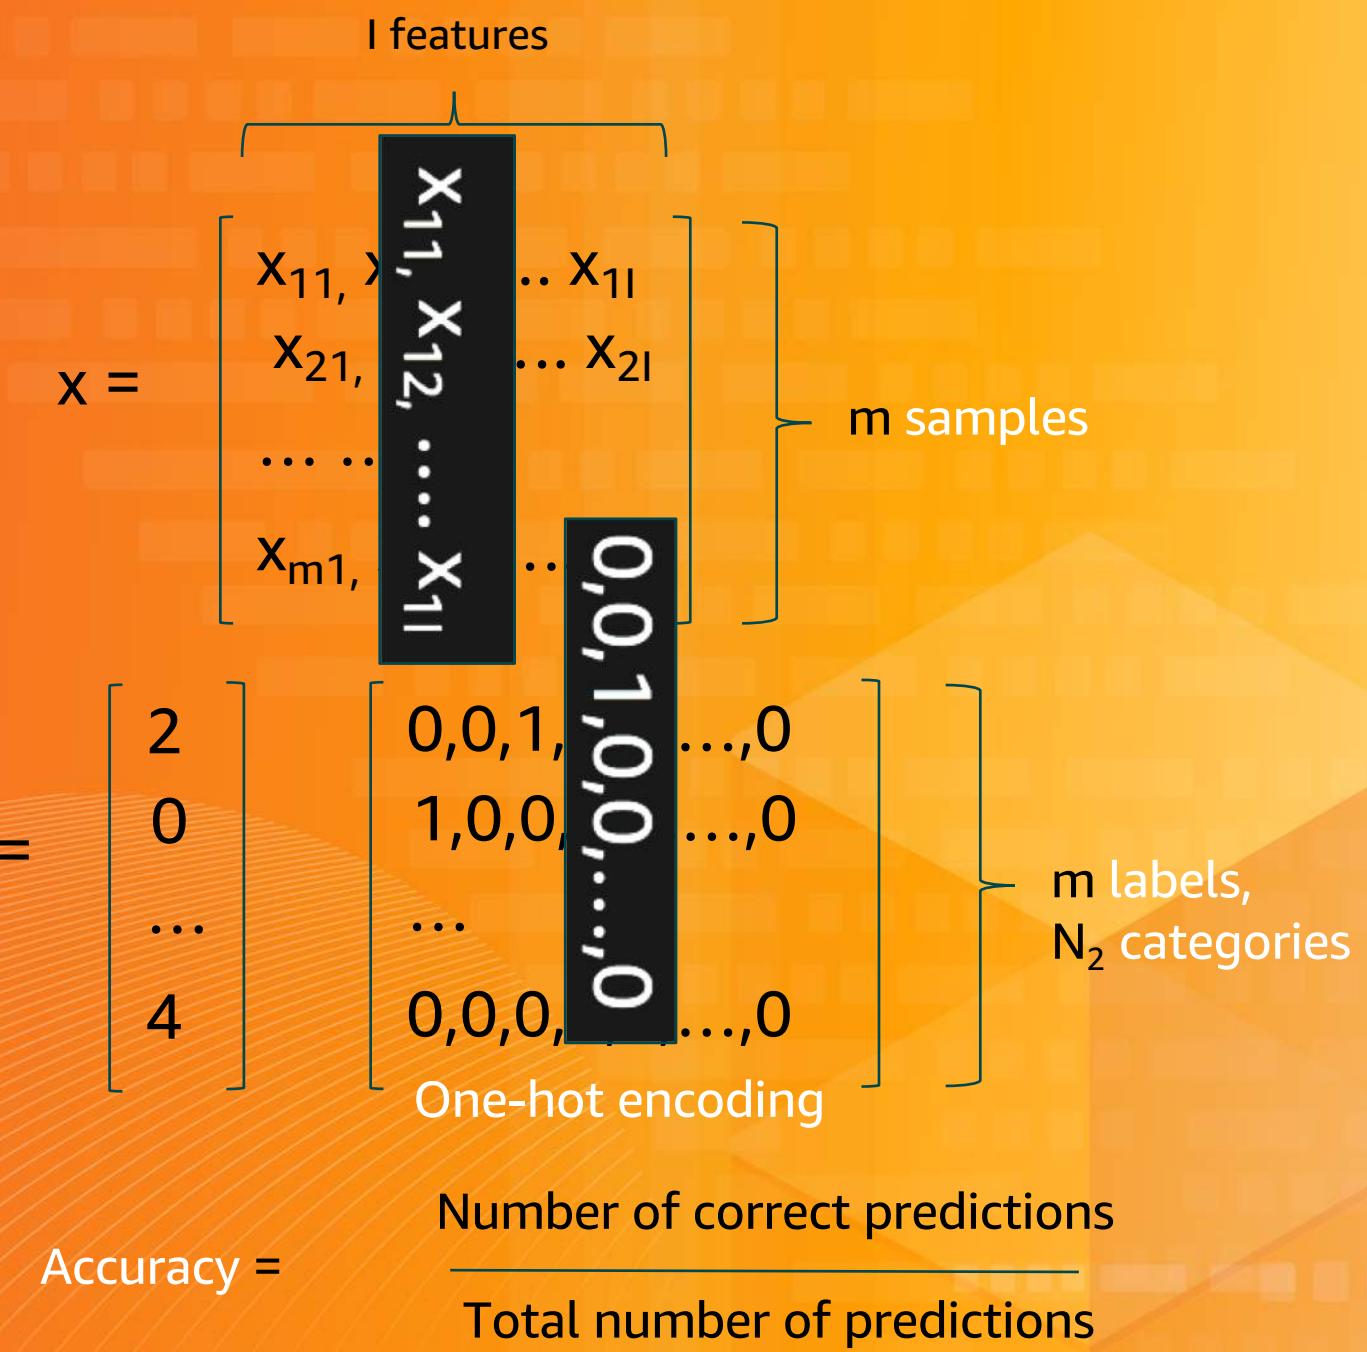
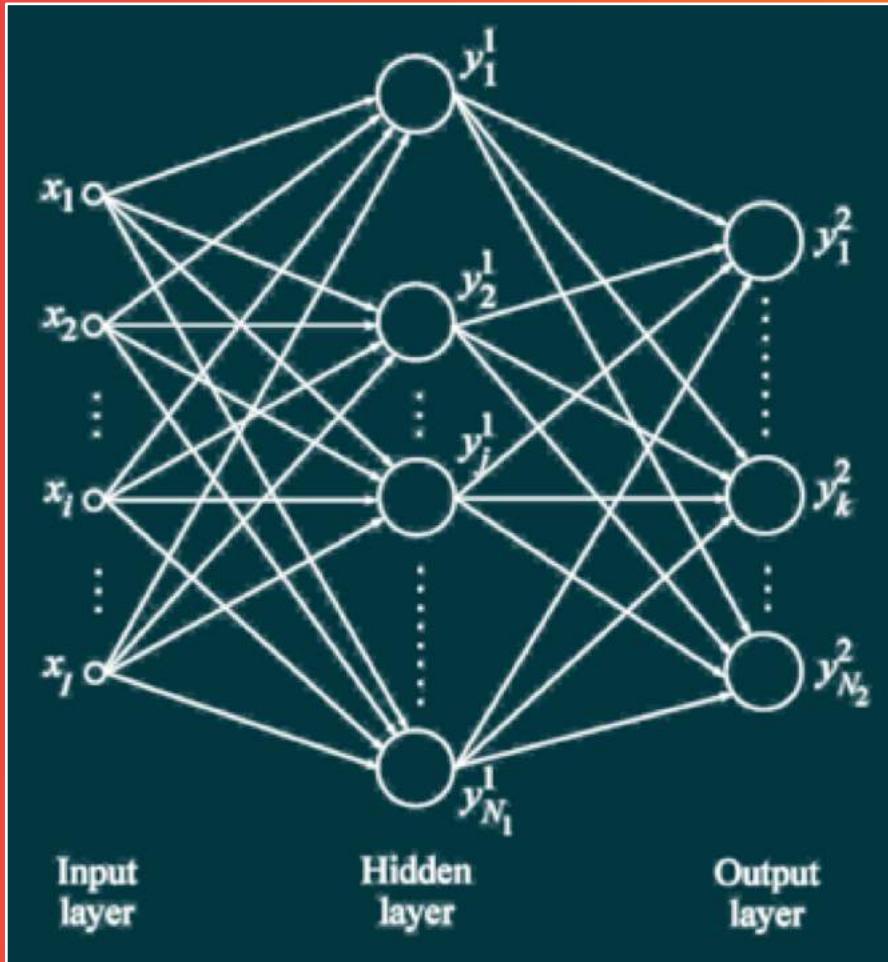
Source: Wikipedia



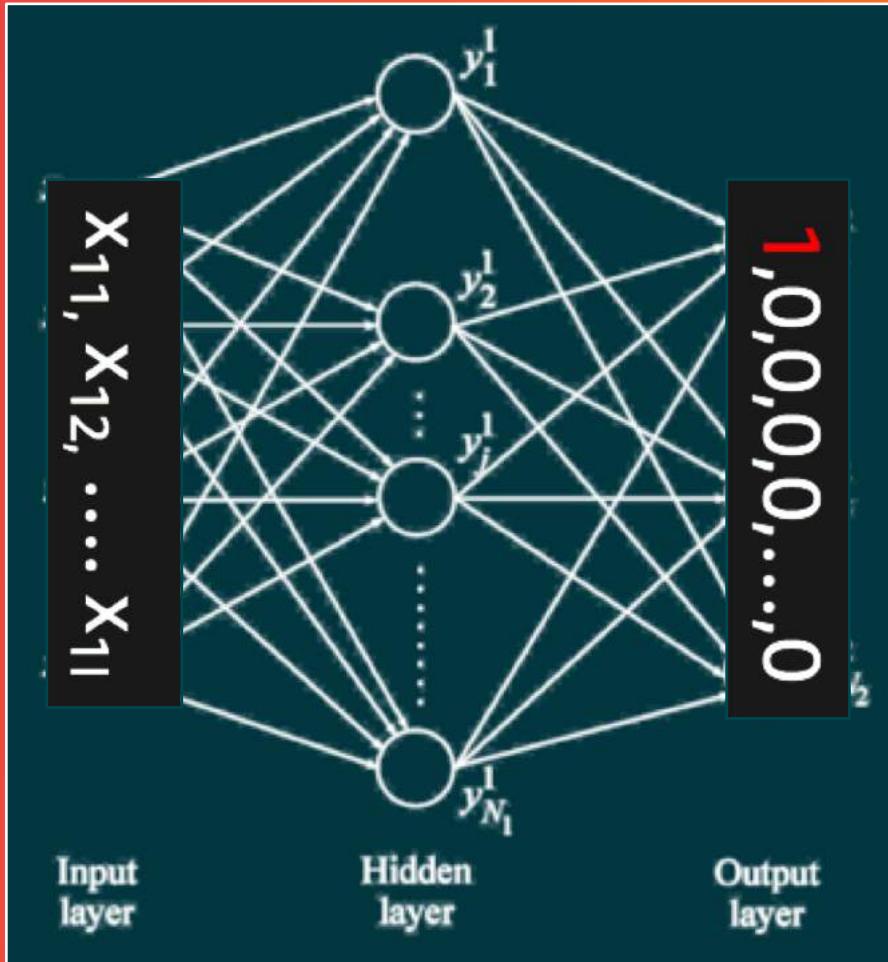
Neural networks



Neural networks



Neural networks



Initially, the network will not predict correctly

$$f(X_1) = Y'_1$$

A loss function measures the difference between the real label Y_1 and the predicted label Y'_1
error = loss(Y_1, Y'_1)

For a batch of samples:

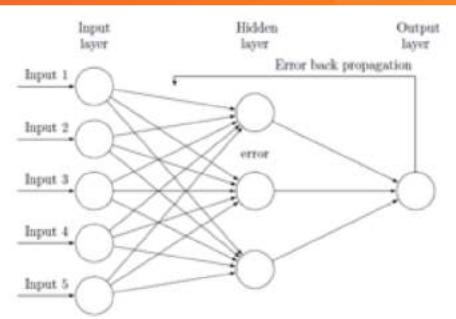
$$\text{batch size} \sum_{i=1}^{batch\ size} \text{loss}(Y_i, Y'_i) = \text{batch error}$$

The purpose of the training process is to minimize error by gradually adjusting weights.

Training



Training data set

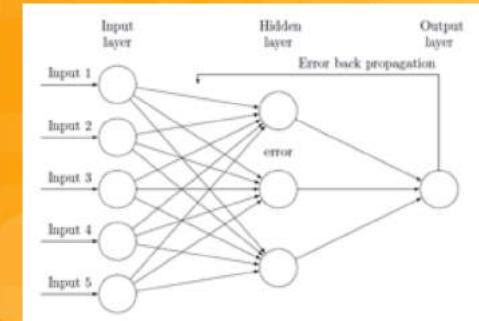


Backpropagation

Batch size
Learning rate
Number of epochs



Hyper parameters

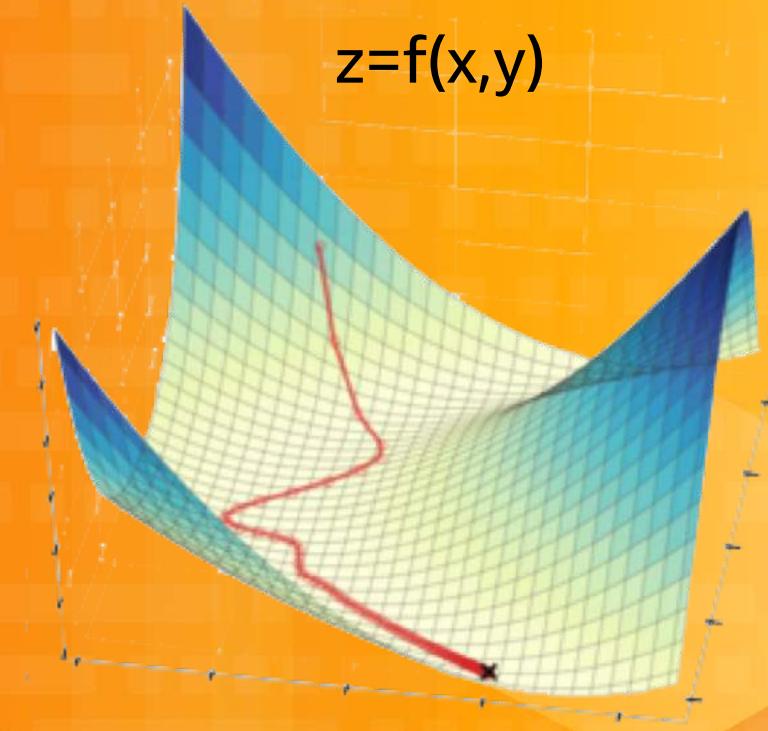


Trained
neural network

Stochastic Gradient Descent (1951)

Imagine you stand on top of a mountain with skis strapped to your feet. You want to get down to the valley as quickly as possible, but there is fog and you can only see your immediate surroundings. How can you get down the mountain as quickly as possible? You look around and identify the steepest path down, go down that path for a bit, again look around and find the new steepest path, go down that path, and repeat—this is exactly what gradient descent does.

Tim Dettmers
University of Lugano
2015



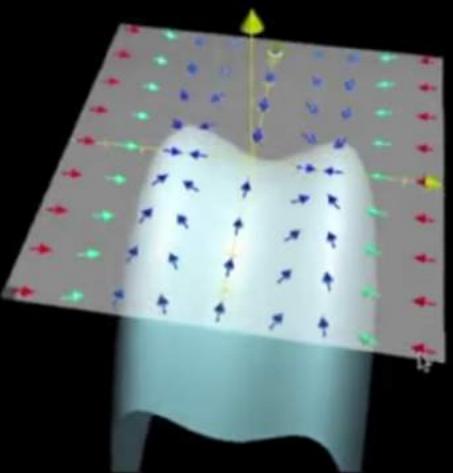
The « step size » depends
on the learning rate

<https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-history-training/>

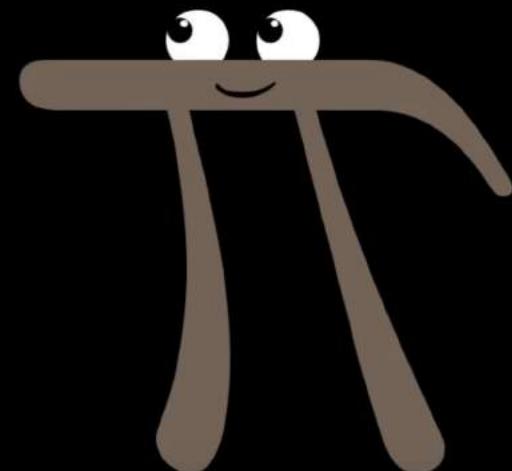
$$f(x, y) = x^2 + y^2$$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

$$= \begin{bmatrix} 2x \\ 2y \end{bmatrix}$$



direction of
steepest
ascent



Error derivative

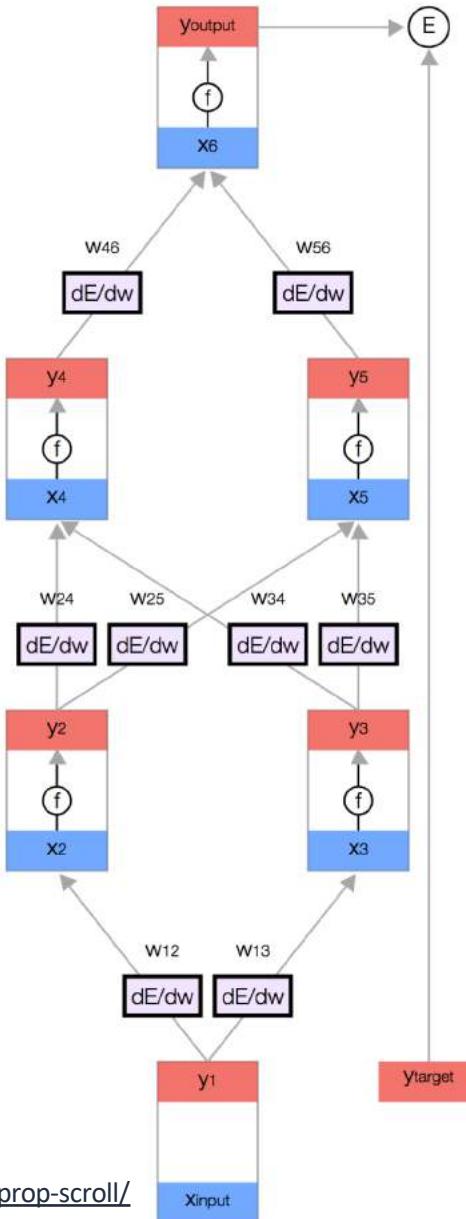
The backpropagation algorithm decides how much to update each weight of the network after comparing the predicted output with the desired output for a particular example. For this, we need to compute how the error changes with respect to each weight $\frac{dE}{dw_{ij}}$.

Once we have the error derivatives, we can update the weights using a simple update rule:

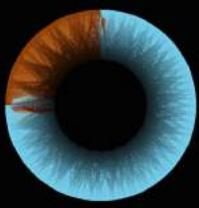
$$w_{ij} = w_{ij} - \alpha \frac{dE}{dw_{ij}}$$

where α is a positive constant, referred to as the learning rate, which we need to fine-tune empirically.

[Note] The update rule is very simple: if the error goes down when the weight increases ($\frac{dE}{dw_{ij}} < 0$), then increase the weight, otherwise if the error goes up when the weight increases ($\frac{dE}{dw_{ij}} > 0$), then decrease the weight.



<https://google-developers.appspot.com/machine-learning/crash-course/backprop-scroll/>



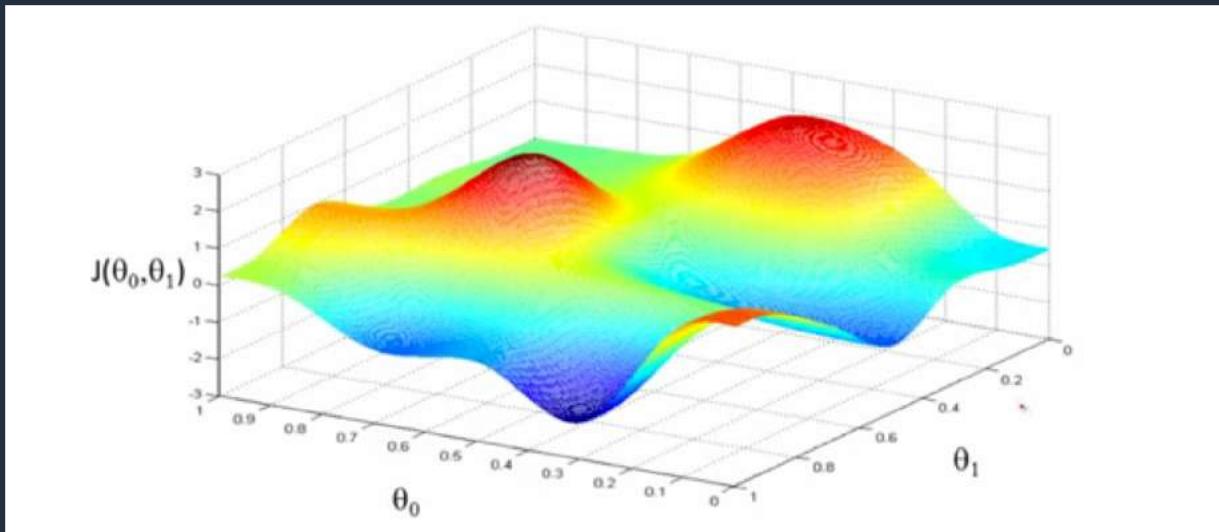
$$\vec{\mathbf{W}} = \begin{bmatrix} 2.25 \\ -1.57 \\ 1.98 \\ \vdots \\ -1.16 \\ 3.82 \\ 1.21 \end{bmatrix} \quad 0.18$$

$$-\nabla C(\vec{\mathbf{W}}) =$$

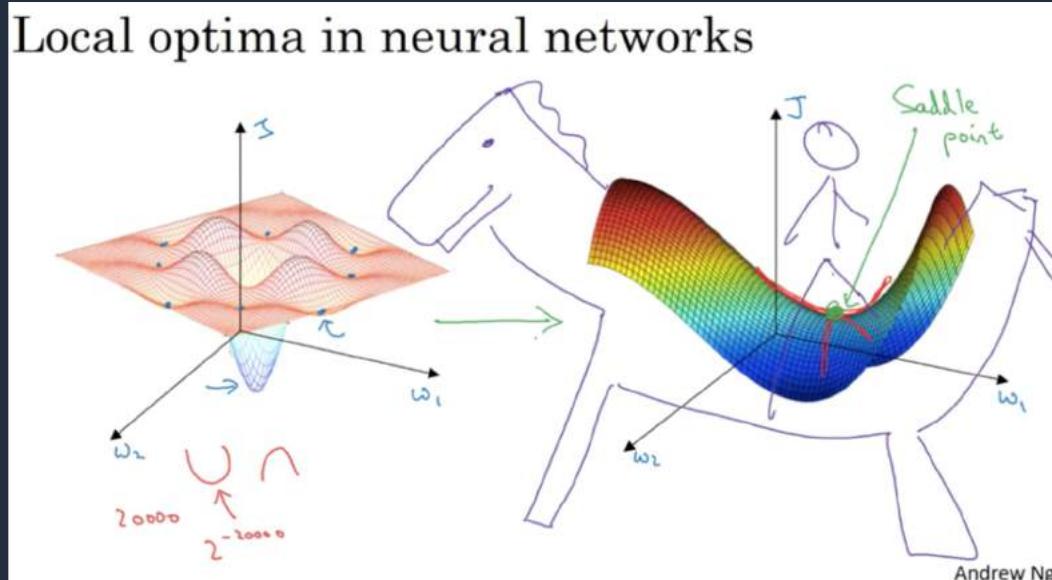
$$\begin{bmatrix} 0.18 \\ 0.45 \\ -0.51 \\ \vdots \\ 0.40 \\ -0.32 \\ 0.82 \end{bmatrix}$$



Local minima and saddle points



Local optima in neural networks



Andrew Ng

« Do neural networks enter and escape a series of local minima? Do they move at varying speed as they approach and then pass a variety of saddle points? Answering these questions definitively is difficult, but we present evidence strongly suggesting that the answer to all of these questions is no. »

« Qualitatively characterizing neural network optimization problems », Goodfellow et al, 2015
<https://arxiv.org/abs/1412.6544>



In neural nets, why use gradient methods rather than other metaheuristics?

For large-size networks, most local minima are equivalent and yield similar performance on a test set.

asked 2 years, 2 months ago

viewed 1,817 times

active 1 year, 9 months ago

In training deep and shallow neural networks, why are gradient methods (e.g. gradient descent, Nesterov, Newton-Raphson) commonly used, as opposed to other metaheuristics?

The probability of finding a "bad" local minimum decreases quickly with networks size.

neural-networks optimization deep-learning gradient-descent backpropagation

share cite edit flag

edited Apr 15 '16 at 7:19

asked Apr 15 '16 at 7:14

5 Silverfish Lior

4:55

Struggling to find the global minimum is not so useful in practice and may lead to overfitting.

FEATURED ON META

Let's hold language in comments to the same standard as posts

HOT META POSTS

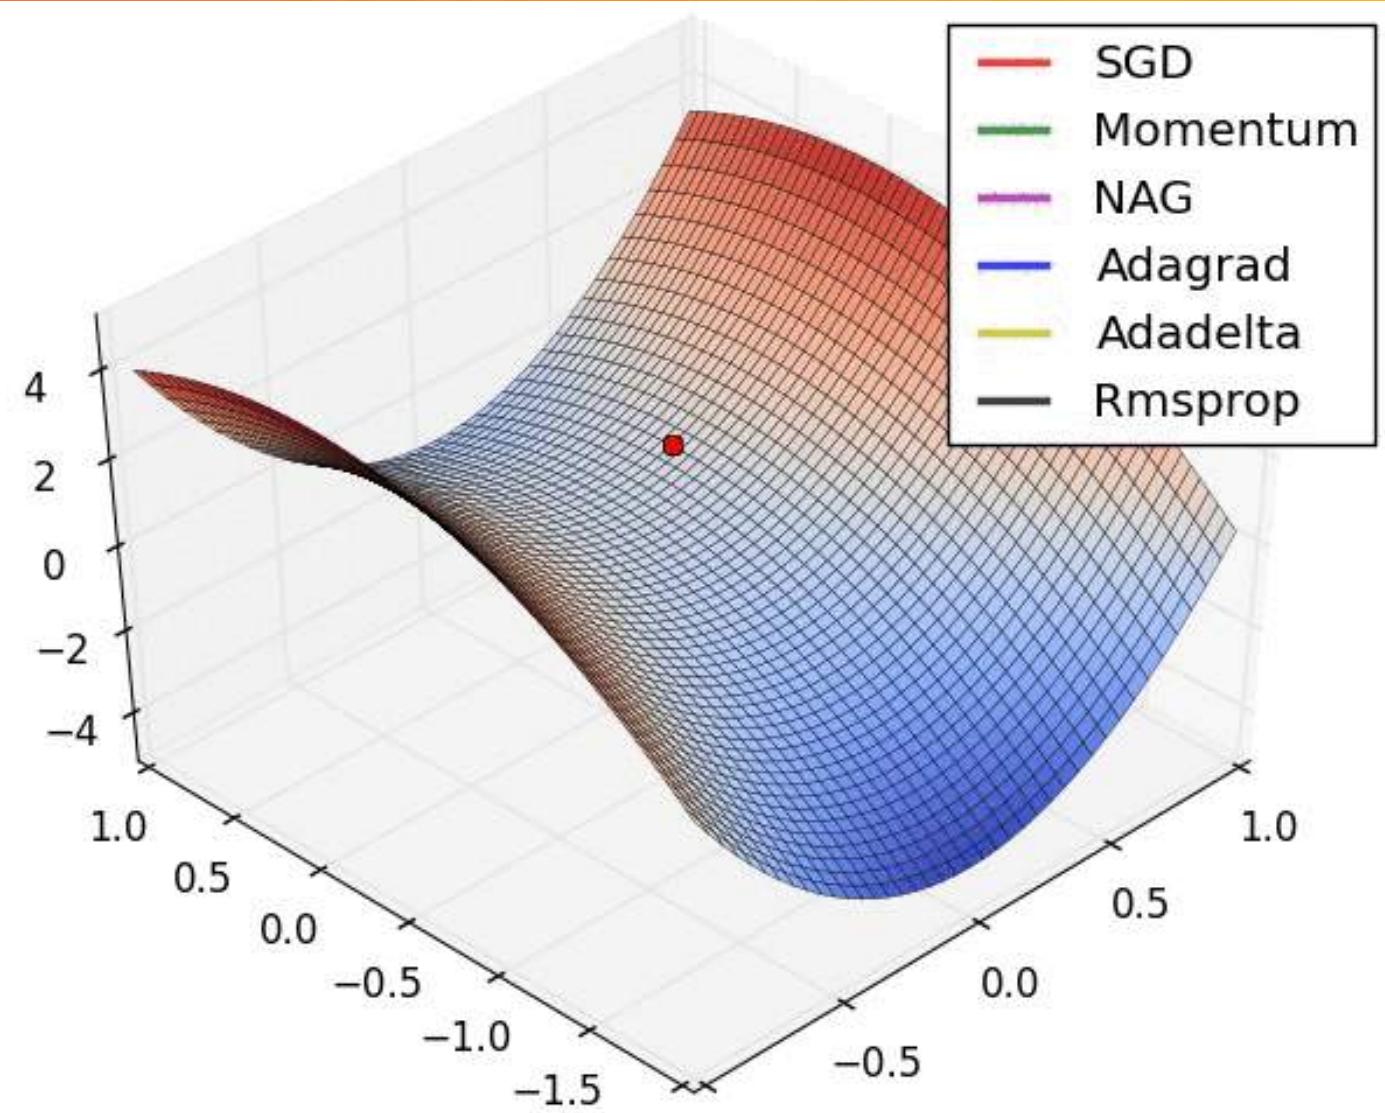
- 6 Need for and name of [acf-pcf] tag
- 7 Why do we have the [tag:coding] tag?
- 7 Should we delete the type-hints tag?
- 12 When we vote to close a question as "unclear" should we state why in a comment?

Optimizers

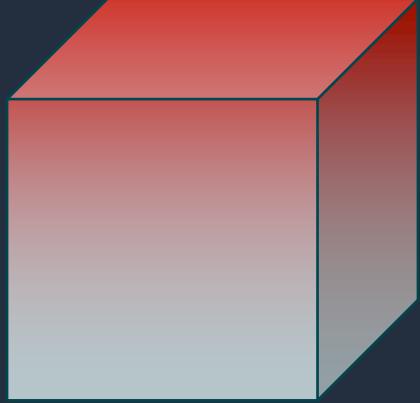
SGD works remarkably well and is still widely used.

Adaptive optimizers use a variable learning rate.

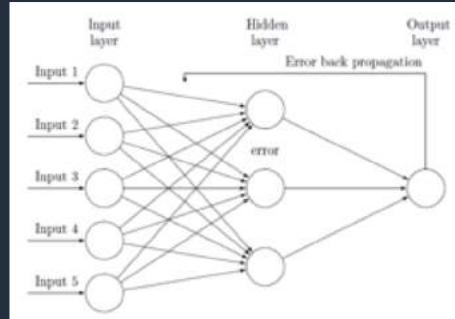
Some even use a learning rate per dimension (Adam).



Validation



Validation data set
(also called dev set)



Neural network
in training



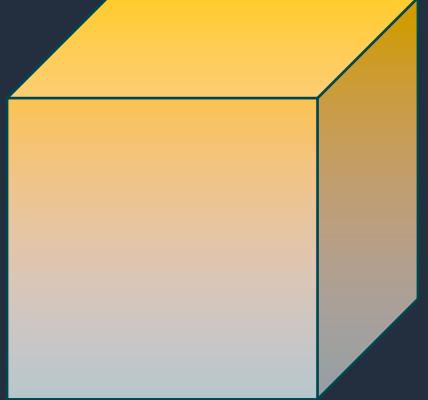
Validation
accuracy

Prediction at
the end of
each epoch

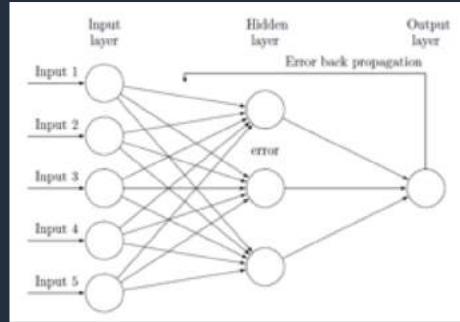
This data set must have the same distribution as real-life samples,
or else validation accuracy won't reflect real-life accuracy.



Test



Test data set



Fully trained
neural network

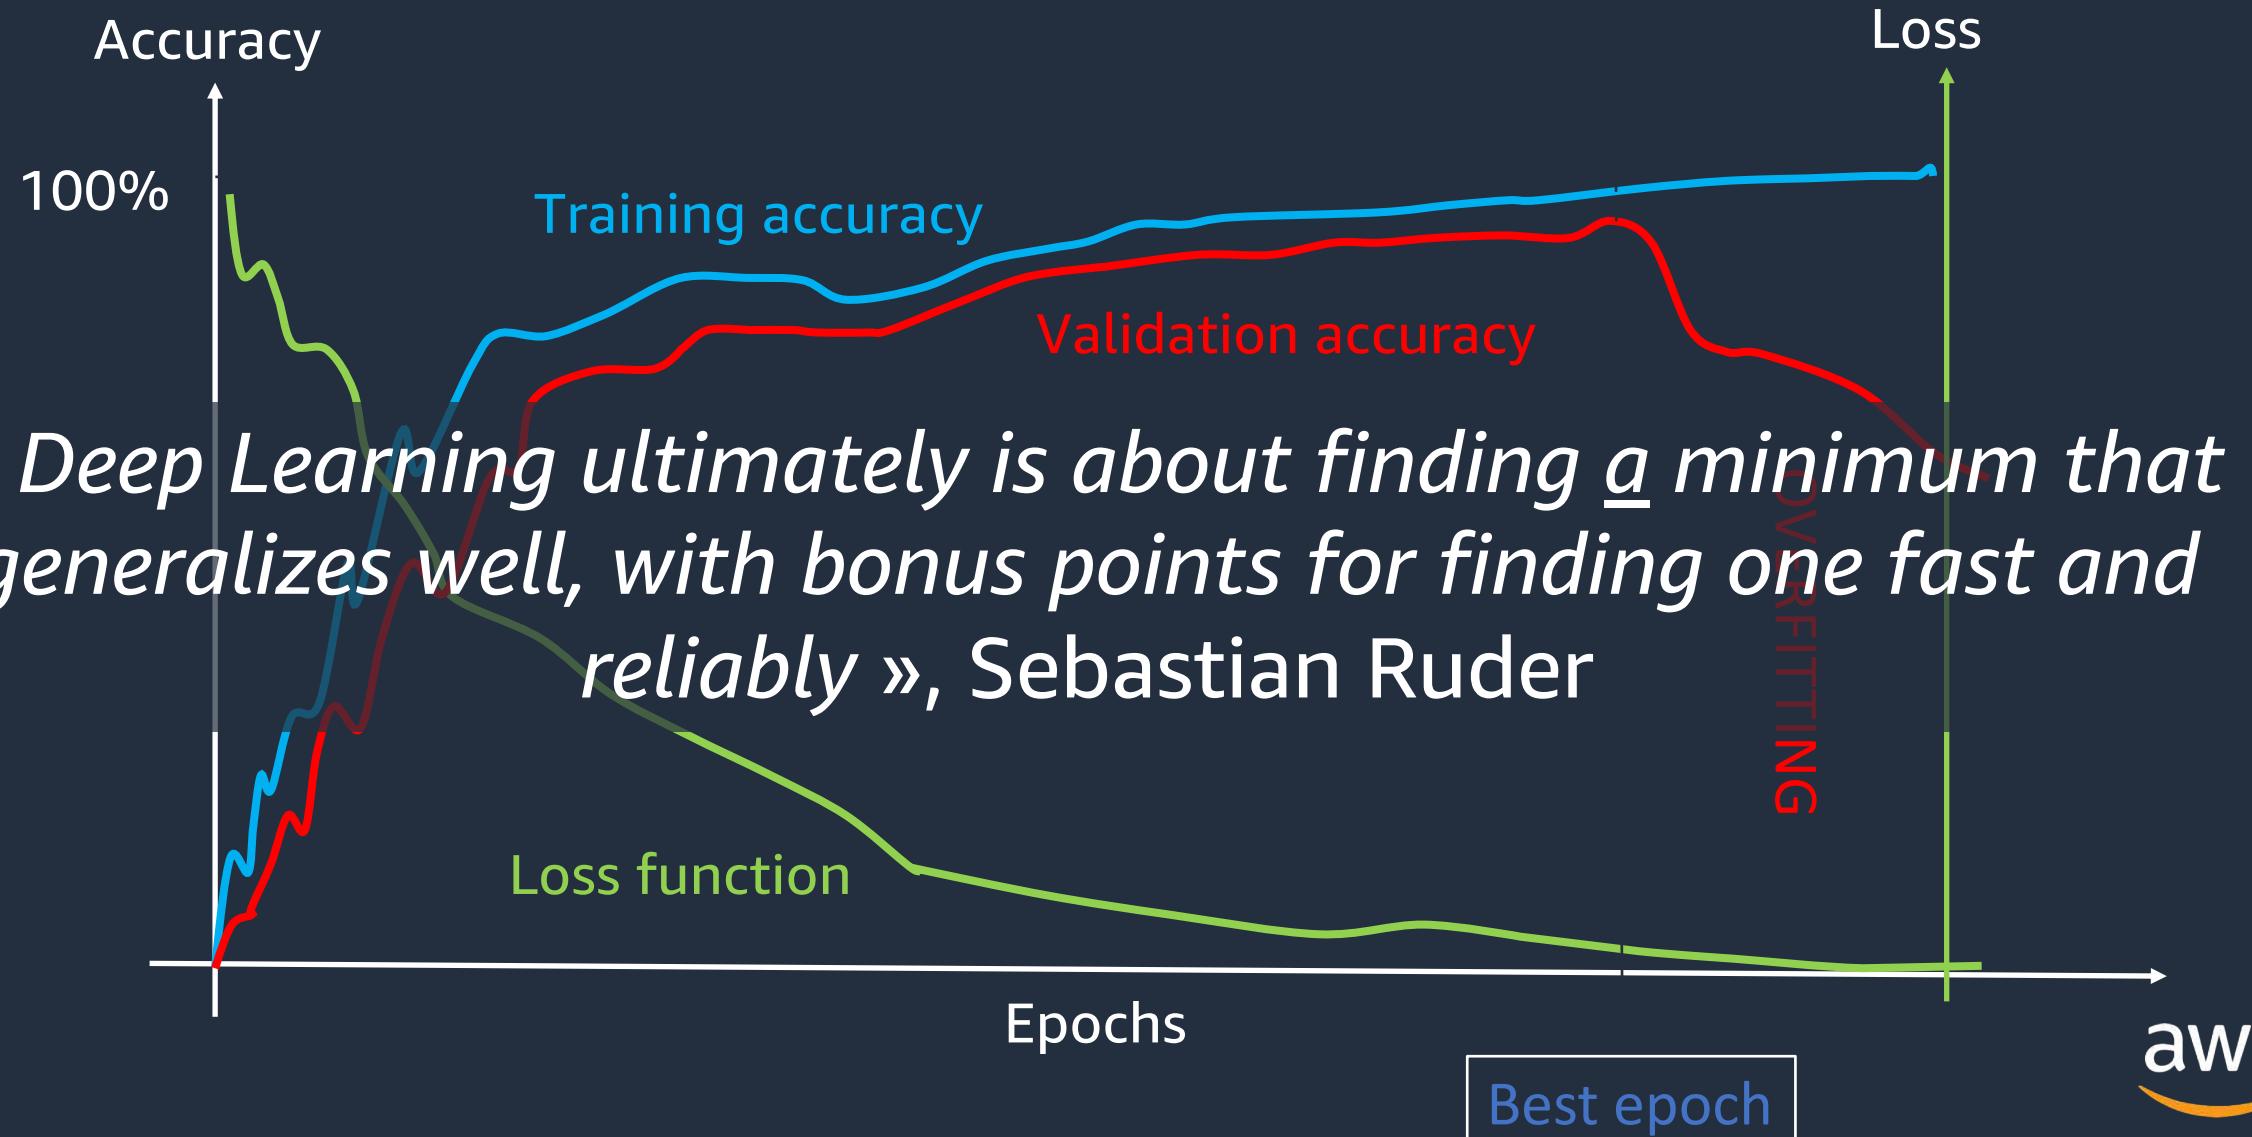


Test accuracy

Prediction at
the end of
experimentation

This data set must have the same distribution as real-life samples,
or else test accuracy won't reflect real-life accuracy.

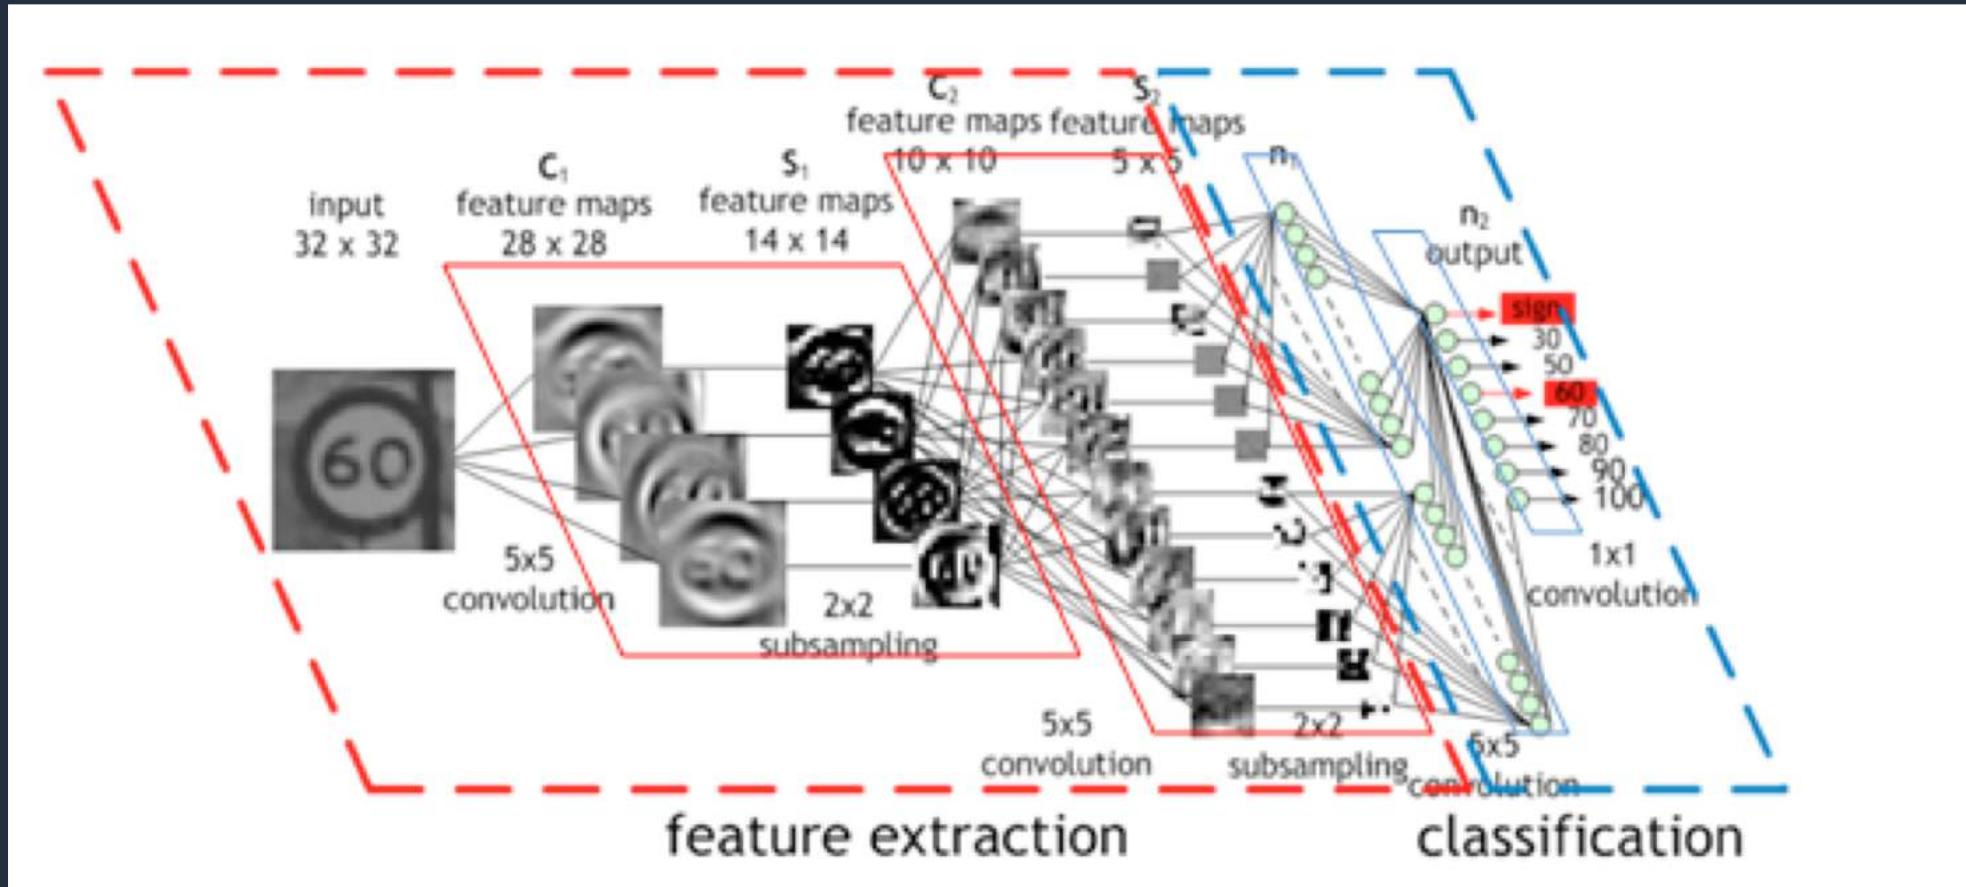
Early stopping



Common network architectures and use cases

Convolutional Neural Networks (CNN)

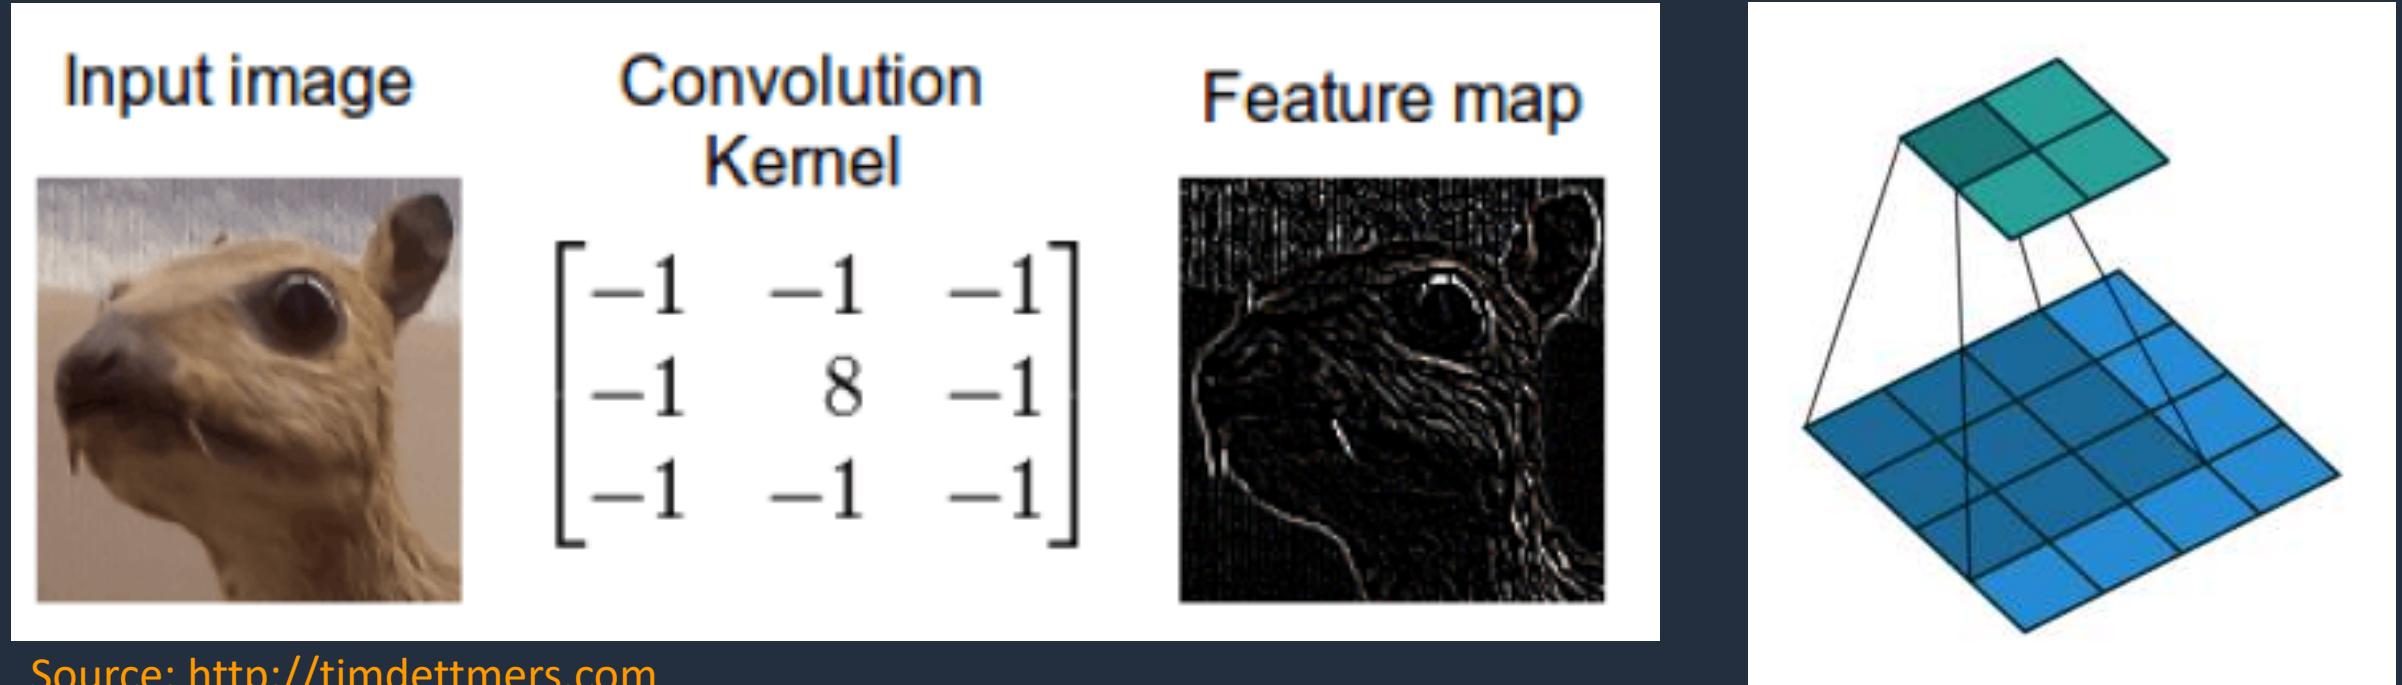
Le Cun, 1998: handwritten digit recognition, 32x32 pixels



<https://devblogs.nvidia.com/parallelforall/deep-learning-nutshell-core-concepts/>



Extracting features with convolution

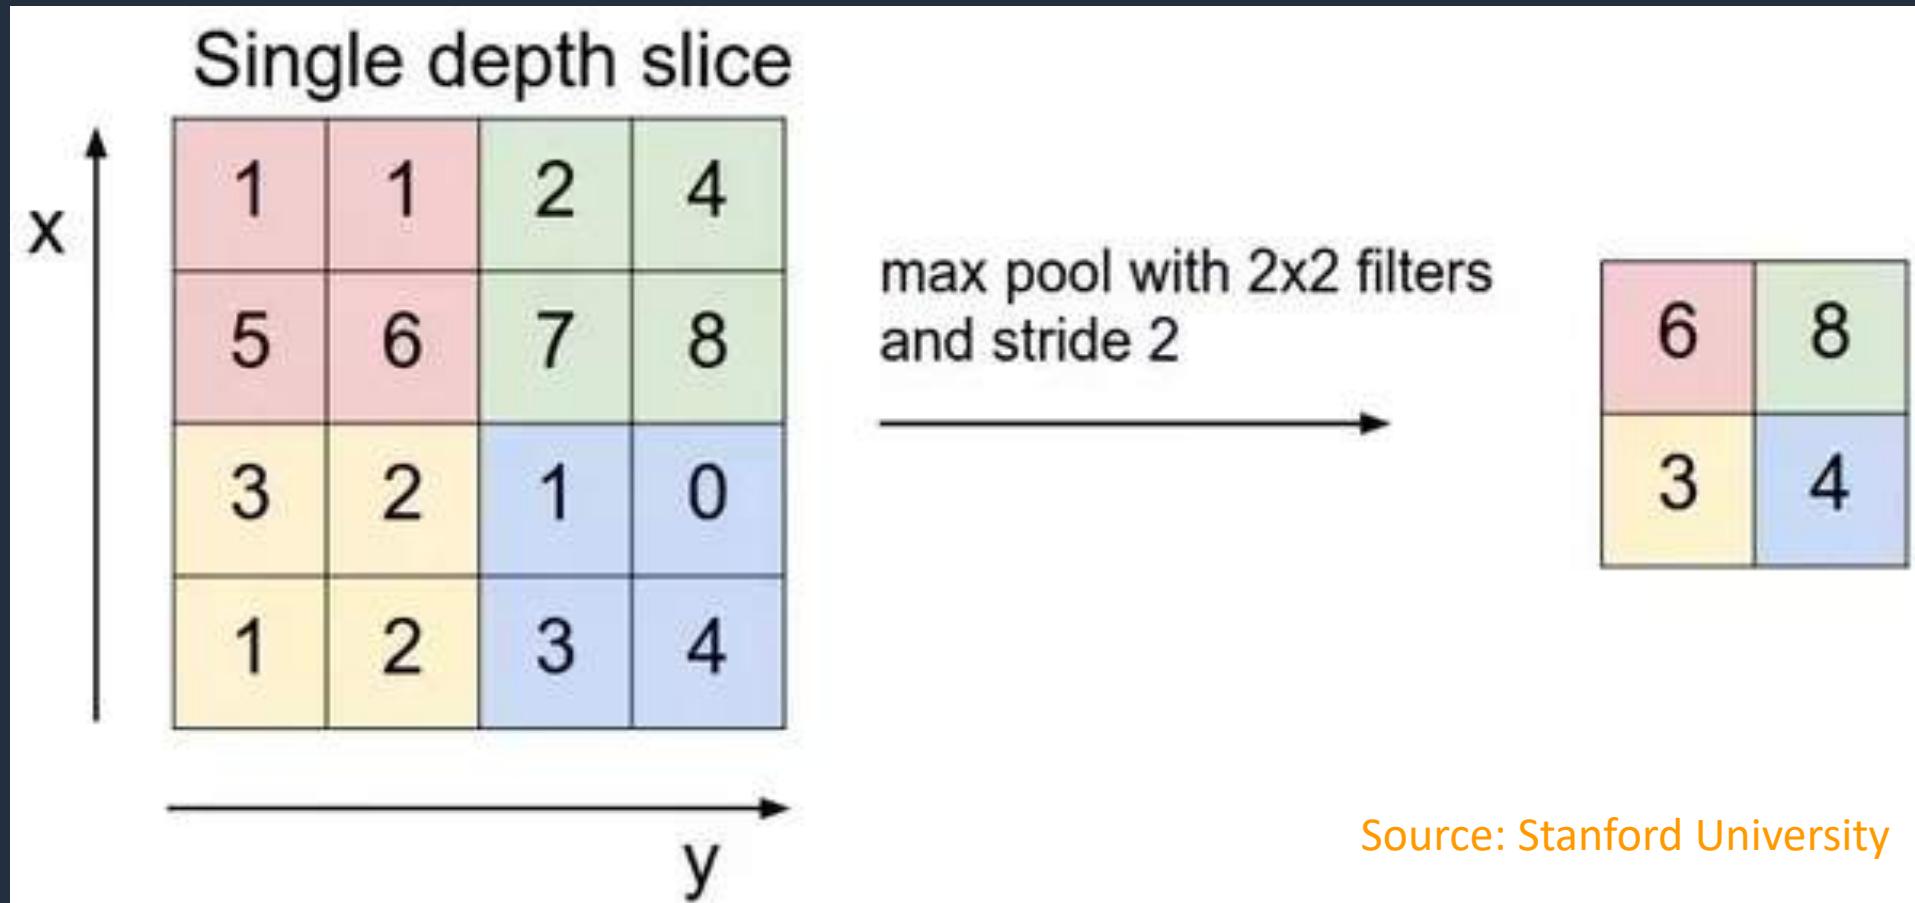


Source: <http://timdettmers.com>

Convolution **extracts features** automatically.
Kernel parameters are **learned** during the training process.



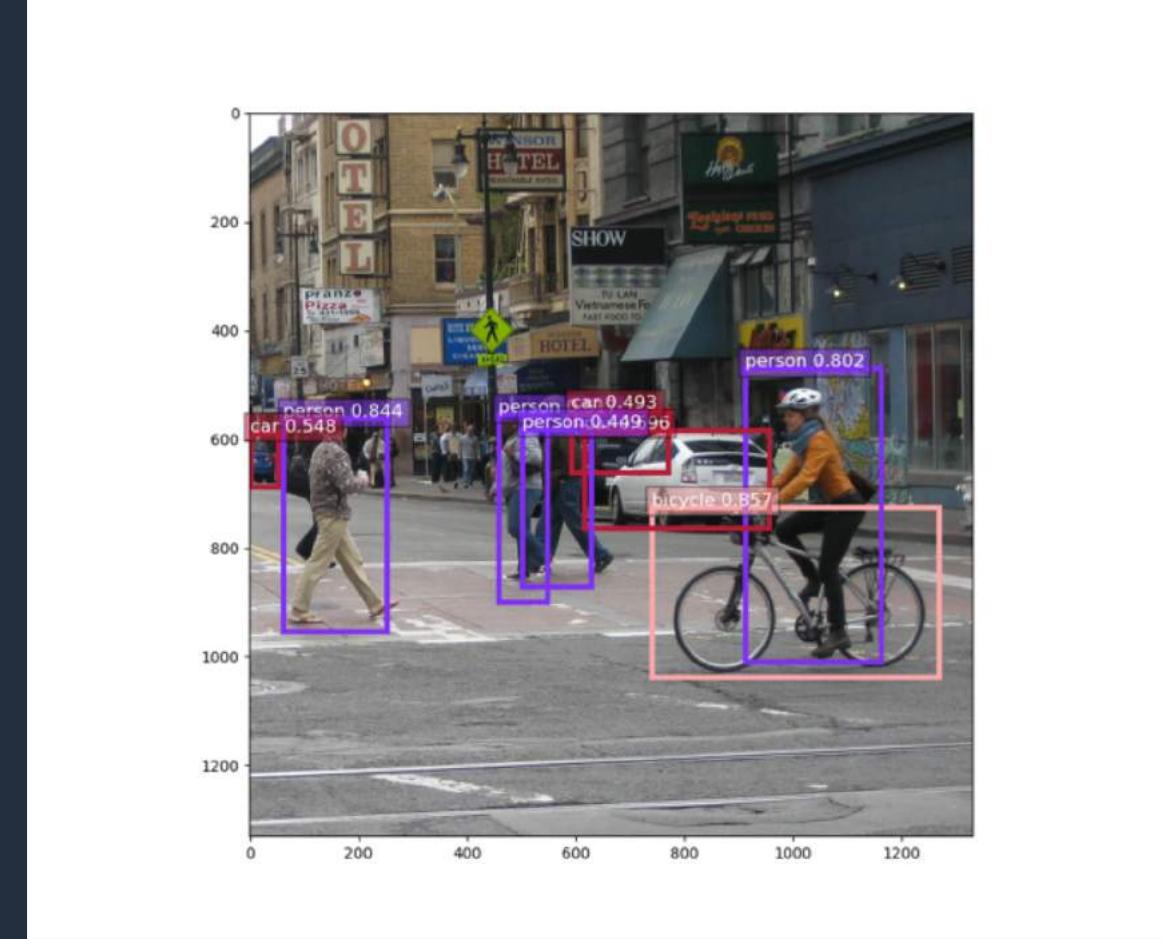
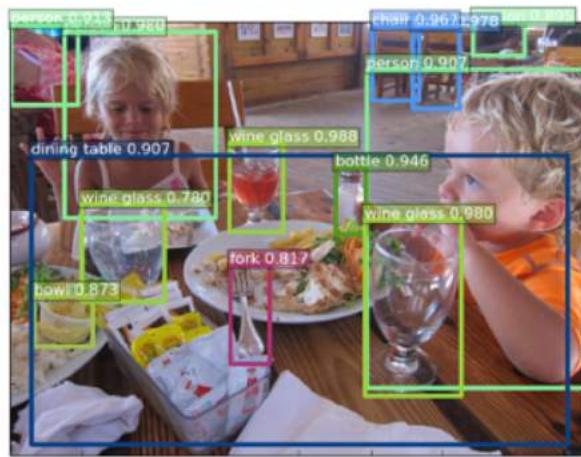
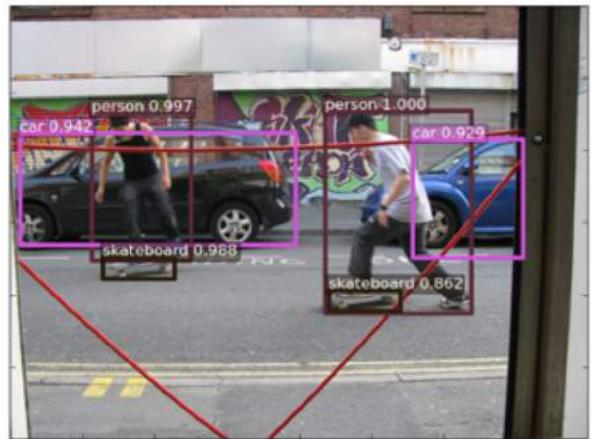
Downsampling images with pooling



Pooling shrinks images while preserving **significant** information.



Object Detection



<https://github.com/precedenceguo/mx-rcnn>

<https://github.com/zhreshold/mxnet-yolo>

Object Segmentation



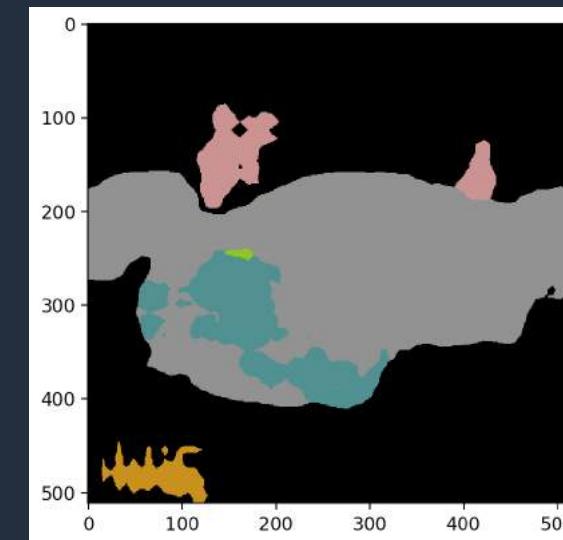
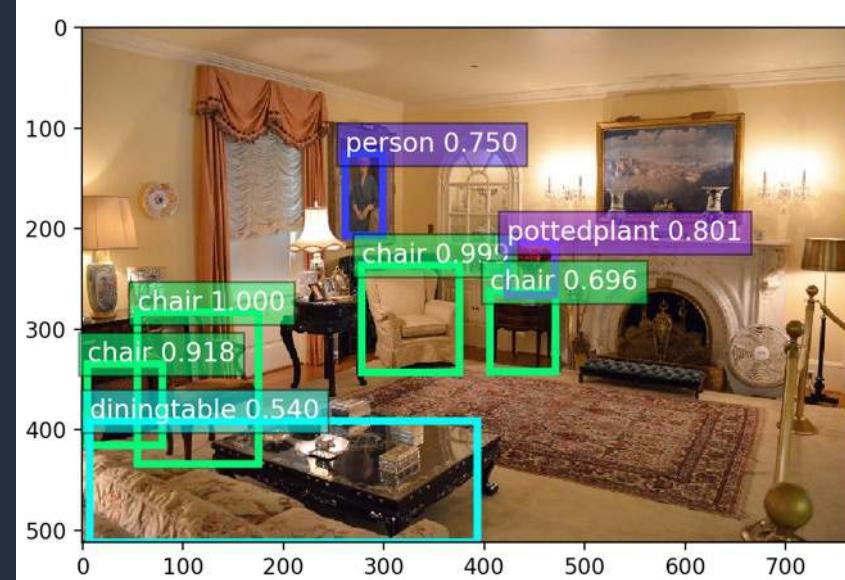
<https://github.com/TuSimple/mx-maskrcnn>



Classification, detection, segmentation



[electric_guitar],
with probability 0.671



Face Detection



<https://github.com/tornadomeet/mxnet-face>

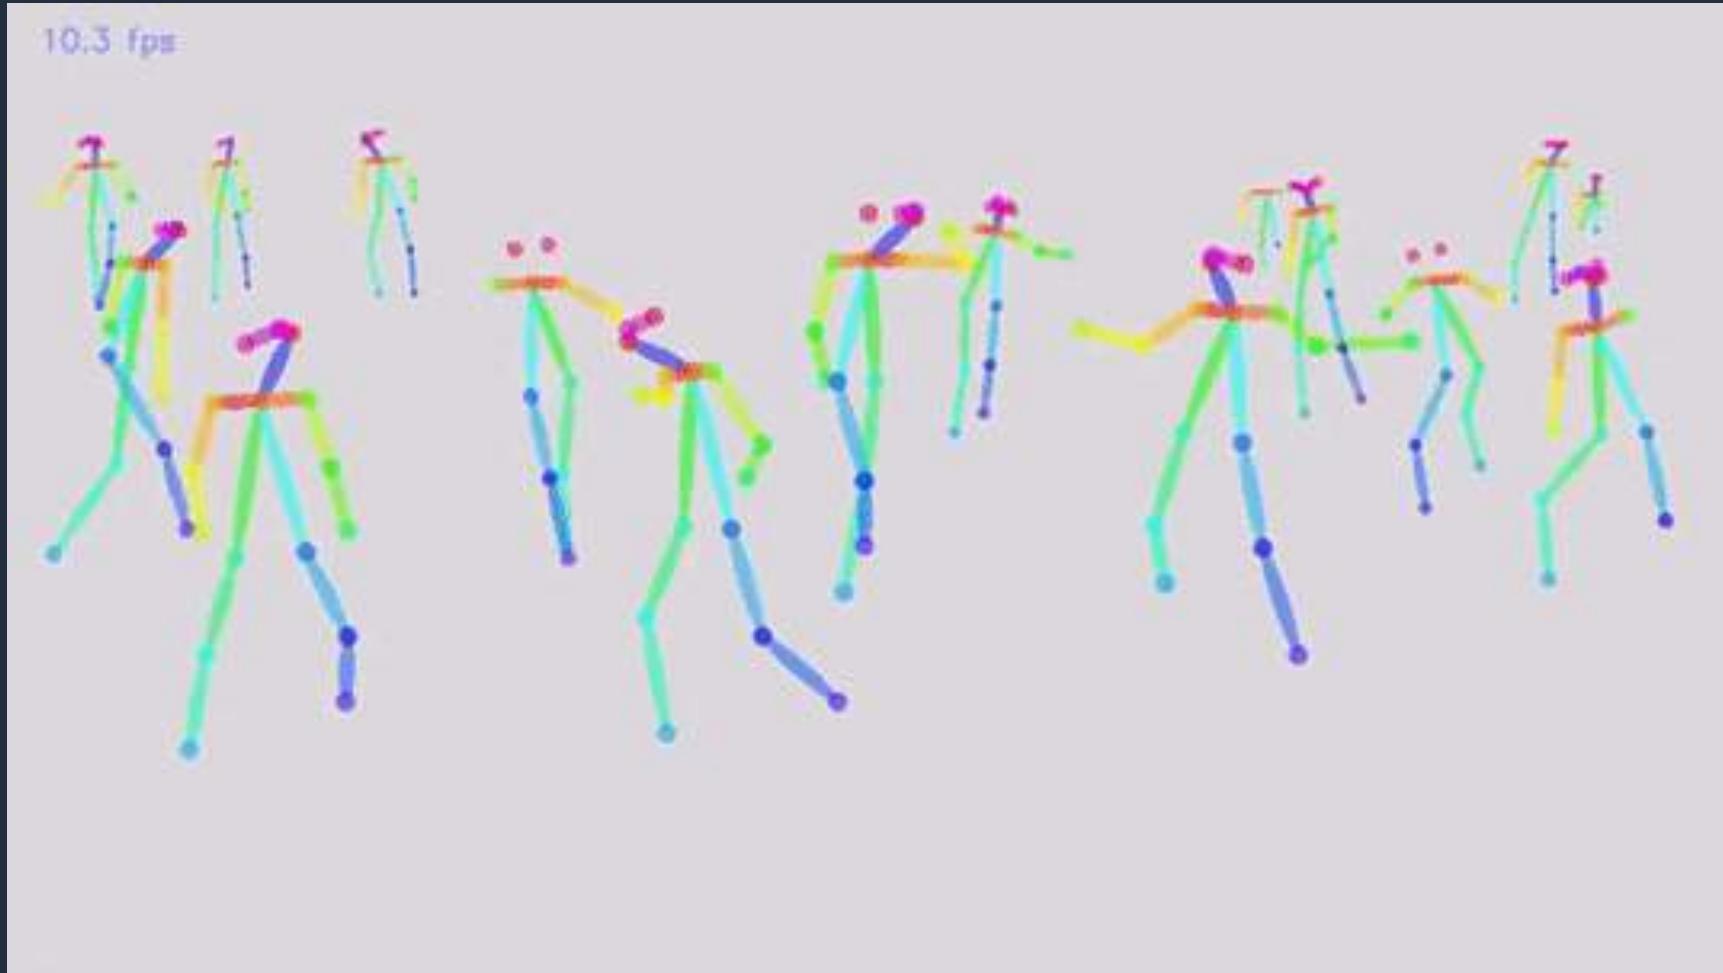
```
attribution is:  
  5_o_Clock_Shadow : No  
  Arched_Eyebrows : No  
  Attractive : No  
  Bags_Under_Eyes : No  
  Bald : No  
  Bangs : No  
  Big_Lips : No  
  Big_Nose : No  
  Black_Hair : No  
  Blond_Hair : No  
  Blurry : Yes  
  Brown_Hair : No  
  Bushy_Eyebrows : No  
  Chubby : No  
  Double_Chin : No  
  Eyeglasses : No  
  Goatee : No  
  Gray_Hair : No  
  Heavy_Makeup : No  
  High_Cheekbones : No  
  Male : Yes  
  Mouth_Slightly_Open : No  
  Mustache : No  
  Narrow_Eyes : Yes  
  No_Beard : Yes  
  Oval_Face : No  
  Pale_Skin : No  
  Pointy_Nose : No  
  Receding_Hairline : No  
  Rosy_Cheeks : No  
  Sideburns : No  
  Smiling : No  
  Straight_Hair : No  
  Wavy_Hair : No  
  Wearing_Earrings : No  
  Wearing_Hat : No  
  Wearing_Lipstick : No  
  Wearing_Necklace : No  
  Wearing_Necktie : No  
  Young : Yes
```

Face Recognition

LFW 99.80%+
Megaface 98%+
with a single model

<https://github.com/deepinsight/insightface>
<https://arxiv.org/abs/1801.07698>

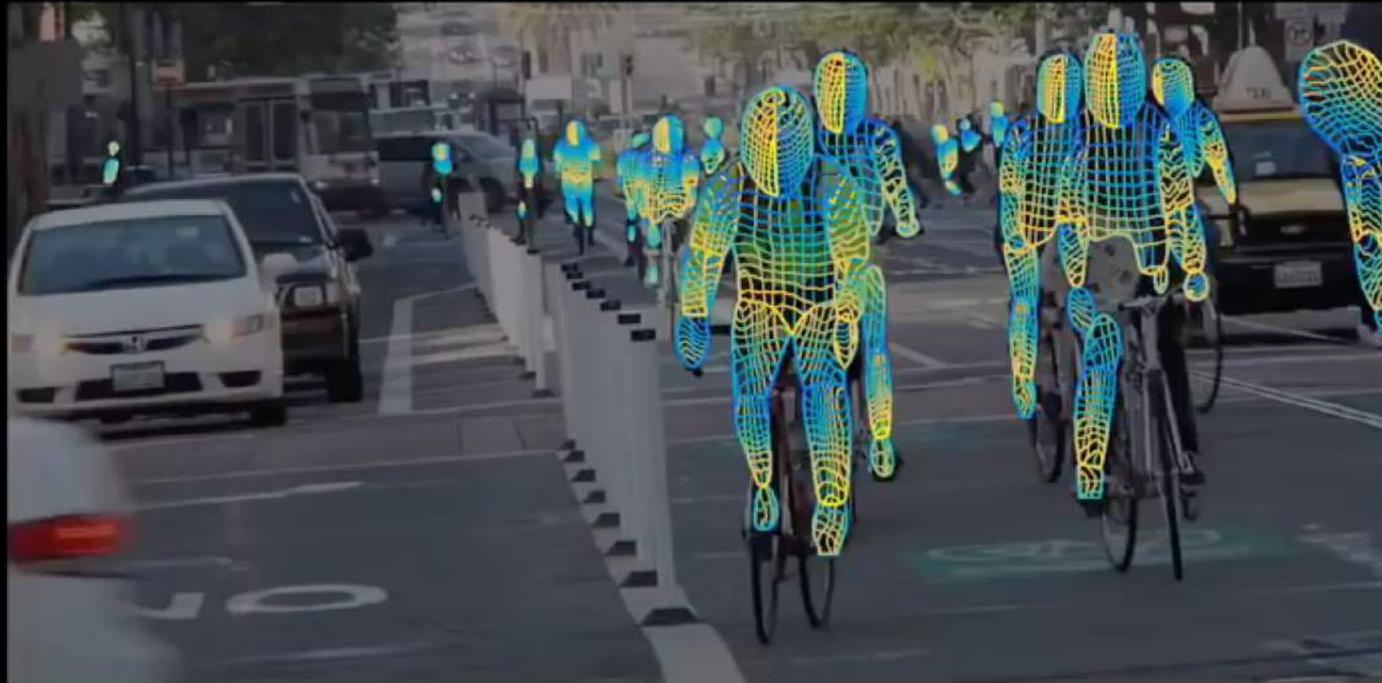
Real-Time Pose Estimation



https://github.com/dragonfly90/mxnet_Realtime_Multi-Person_Pose_Estimation

DensePose:

Dense Human Pose Estimation In The Wild



Rıza Alp Güler *

INRIA, CentraleSupélec

Natalia Neverova

Facebook AI Research

Iasonas Kokkinos

Facebook AI Research

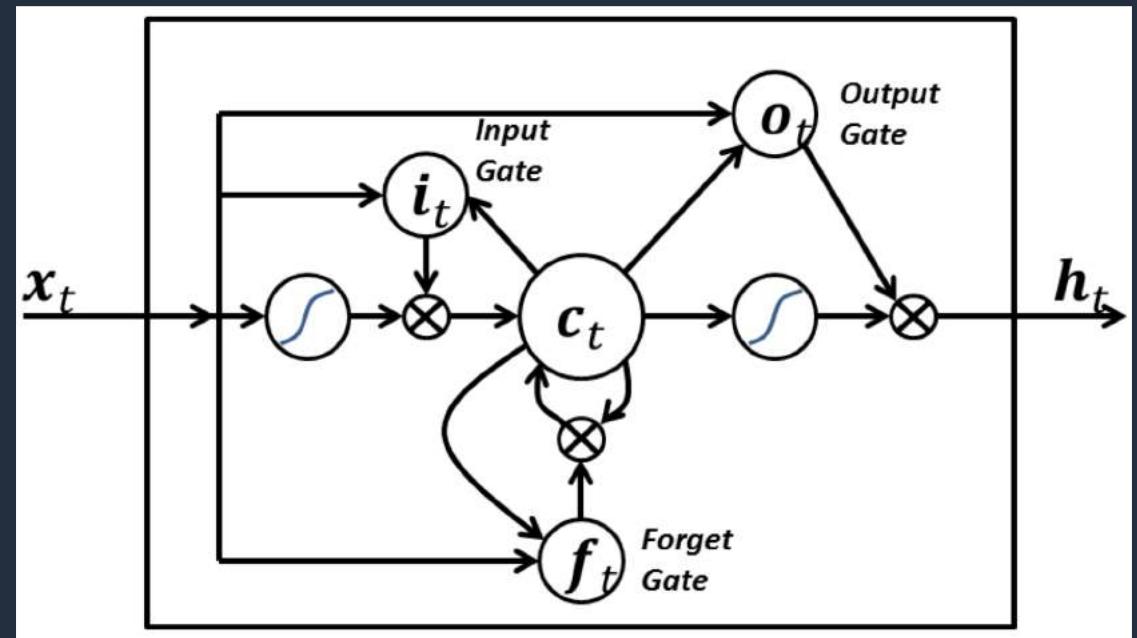
* Rıza Alp Güler was with Facebook AI Research during this work.

Long Short Term Memory Networks (LSTM)

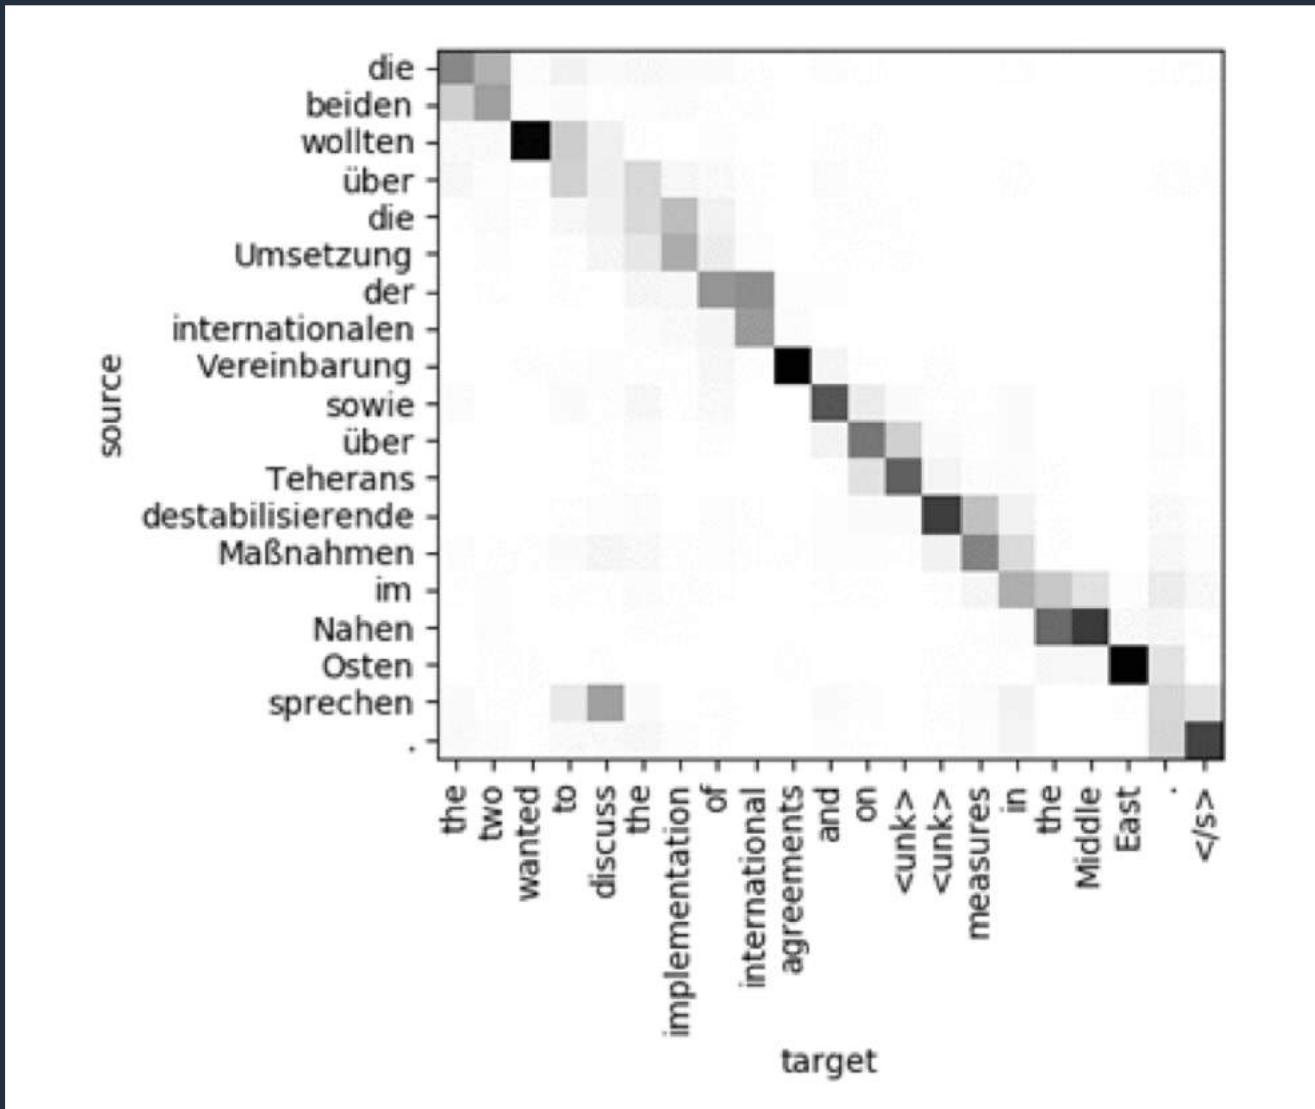
A LSTM neuron computes the output based on the input and a **previous state**

LSTM networks have **memory**

They're great at predicting **sequences** of data



Machine Translation – AWS Sockeye



<https://github.com/awslabs/sockeye>

OCR – Tesseract 4.0 (beta)

Store #05666
3515 DEL MAR HTS,RD
SAN DIEGO, CA 92130
(858) 792-7040

Register #4 Transaction #571140
Cashier #56661020 8/20/17 5:45PM

wellness+ with Plenti
Plenti Card#: 31XXXXXXXXXXXX4553
1 G2 RETRACT BOLD BLK 2PK 1.99 T
SALE 1/1.99, Reg 1/4.69
Discount 2.70-

1 Items	Subtotal	1.99
	Tax	.15
	Total	2.14
MASTER		2.14
MASTER card * #XXXXXXXXXXXX5485		
App #AA APPROVAL AUTO		
Ref # 05639E		
Entry Method: Chip		

OCR Receipt Example

Output

Store #056663515
DEL MAR HTS,RD
SAN DIEGO, CA 92130
(858) 792-7040 Register #4 Transaction
#571140
Cashier #56661020 8/20/17
5:45PM wellness+ with Plenti
Plenti Card#: 31XXXXXXXXXXXX4553
1 G2 RETRACT BOLD BLK 2PK 1.99 T
SALE 1/1.99, Reg 1/4.69
Discount 2.70-

1 Items Subtotal 1.99
Tax .15

Total 2.14
MASTER 2.14
MASTER card * #XXXXXXXXXXXX548S
Apo #AA APPROVAL AUTO
Ref # 05639E
Entry Method: Chip

<https://github.com/tesseract-ocr/tesseract/wiki/4.0-with-LSTM>
<https://www.learnopencv.com/deep-learning-based-text-recognition-ocr-using-tesseract-and-opencv/>



Generative Adversarial Networks

TF



PyTorch



Generating new "celebrity" faces

https://github.com/tkarras/progressive_growing_of_gans

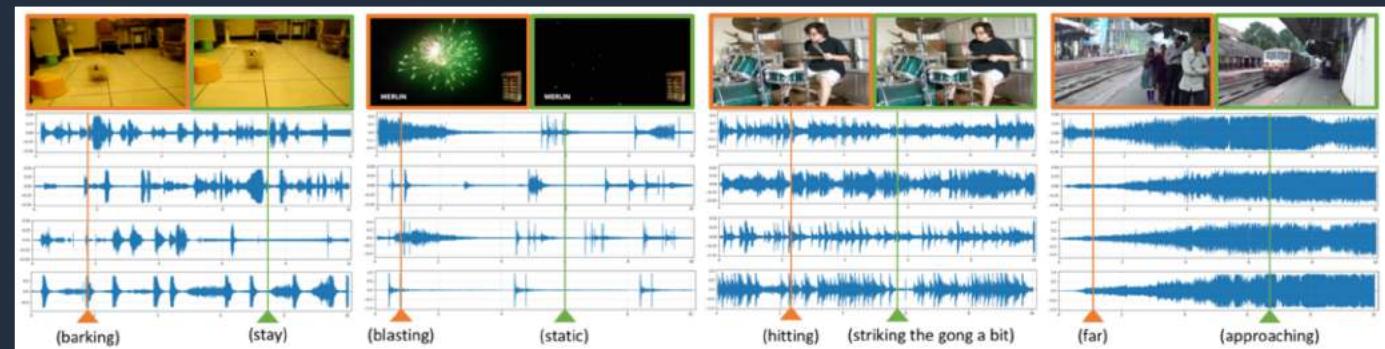
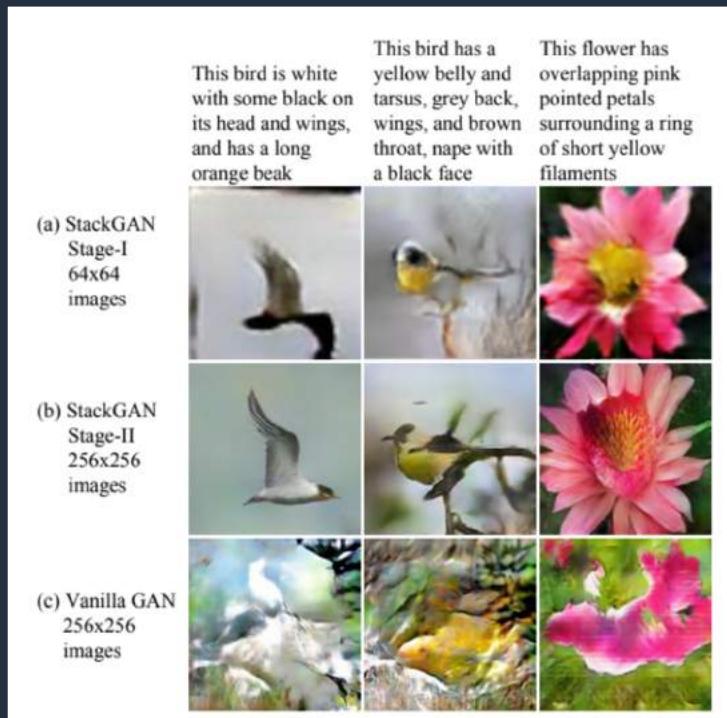
From semantic map to 2048x1024 picture

<https://tcwang0509.github.io/pix2pixHD/>



Wait! There's more!

Models can also generate text from text, text from images, text from video, images from text, sound from video, 3D models from 2D images, etc.





<https://aws.amazon.com/solutions/case-studies/irobot/>

@pitbullsharky

Applause from Adrian Hornsby and 33 others



Julien Simon

Hacker. Headbanger. Harley rider. Hunter. <https://aws.amazon.com/evangelists/julien-simon/>

Jan 12 · 5 min read

10 steps on the road to Deep Learning (part 1)

One of the questions I often get after my talks is: “*I’m a developer. How can I get started with this stuff?*”. Here’s how I worked my way into Deep Learning. By no means am I claiming that this is “The Way”, if such a thing even exists.

In this post and the next, I’ll go through **10 steps** that will hopefully help you learn in the right order and at your own pace.



0.1

[Search docs](#)

CRASH COURSE

[Preface](#)[Introduction](#)[Manipulate data the MXNet way with
ndarray](#)[Linear algebra](#)[Intermediate linear algebra](#)[Probability and statistics](#)[Automatic differentiation with autograd](#)

INTRODUCTION TO SUPERVISED LEARNING

[Linear regression from scratch](#)[Linear regression with gluon](#)[Binary classification with logistic](#)

Deep Learning - The Straight Dope

This repo contains an incremental sequence of notebooks designed to teach deep learning, [Apache MXNet \(incubating\)](#), and the gluon interface. Our goal is to leverage the strengths of Jupyter notebooks to present prose, graphics, equations, and code together in one place. If we're successful, the result will be a resource that could be simultaneously a book, course material, a prop for live tutorials, and a resource for plagiarising (with our blessing) useful code. To our knowledge there's no source out there that teaches either (1) the full breadth of concepts in modern deep learning or (2) interleaves an engaging textbook with runnable code. We'll find out by the end of this venture whether or not that void exists for a good reason.

Another unique aspect of this book is its authorship process. We are developing this resource fully in the public view and are making it available for free in its entirety. While the book has a few primary authors to set the tone and shape the content, we welcome contributions from the community and hope to coauthor chapters and entire sections with experts and community members. Already we've received contributions spanning typo corrections through full working examples.

How to contribute

To clone or contribute, visit [Deep Learning - The Straight Dope](#) on Github.

Getting started

<https://aws.amazon.com/machine-learning>

<https://aws.amazon.com/blogs/ai>

<https://mxnet.incubator.apache.org> | <https://github.com/apache/incubator-mxnet>

<https://gluon.mxnet.io> | <https://github.com/gluon-api>

<https://medium.com/@julsimon>

<https://youtube.com/juliensimonfr>

<https://gitlab.com/juliensimon/dlnotebooks>



AWS DeepLens



aws



INNOVATE2018

ONLINE CONFERENCE

DEVELOPER EDITION

Thank you!

Julien Simon @julsimon

Julio Faerman @faermanj

@awscloud



#AWSInnovate