



# INNOVATE2018

## ONLINE CONFERENCE

DEVELOPER EDITION

# Data Design and Modeling for Microservices

**John Streit, Sr Mgr Specialist Solution Architecture – Data & Analytics, EMEA | 2018**

[jostreit@amazon.com](mailto:jostreit@amazon.com)

@awscloud



#AWSInnovate

# What to Expect from the Session

- Microservices at Amazon
  - Overview and Challenges
  - Key Elements and Benefits
  - Two Pizza Teams
- Data Architecture Challenges
  - Transactions and Rollbacks
  - Streams
  - Master Data Management
- Choosing a Data Store
- Aggregation



# Microservices at Amazon

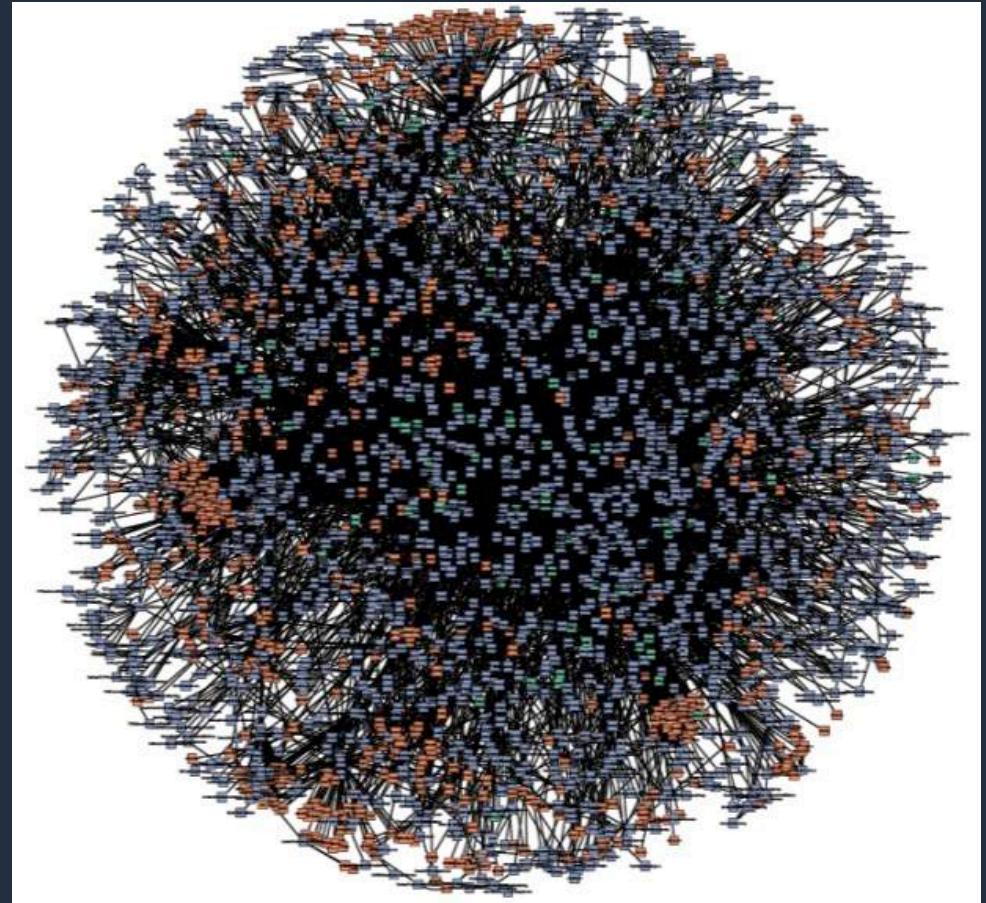
Service-Oriented Architecture (SOA)

Single-purpose

Connect only through APIs

Connect over HTTPS

"Microservices"



# Monolithic vs. SOA vs. Microservices



**Monolithic**  
Single Unit

**SOA**  
Coarse-grained

**Microservices**  
Fine-grained



# Monolithic vs. SOA vs. Microservices

## Microservices:

Many very small components

Business logic lives inside of single service domain

Simple wire protocols(HTTP with XML/JSON)

API driven with SDKs/Clients

## SOA:

Fewer more sophisticated components

Business logic can live across domains

Enterprise Service Bus like layers between services

Middleware



# Microservice Challenges

Distributed computing is hard

Transactions

- Multiple Databases across multiple services

Eventual Consistency

Lots of moving parts

Service discovery

Increase coordination

Increase message routing

# Key Elements of Microservices...

Some core concepts are common to all services

- Service registration, discovery, wiring, administration
- State management
- Service metadata
- Service versioning
- Caching

**Low Friction** Deployment

Automated Management and Monitoring



# Key Elements of Microservices...

Eliminates any long-term commitment to a technology stack

Polyglot **ecosystem**

Polyglot **persistence**

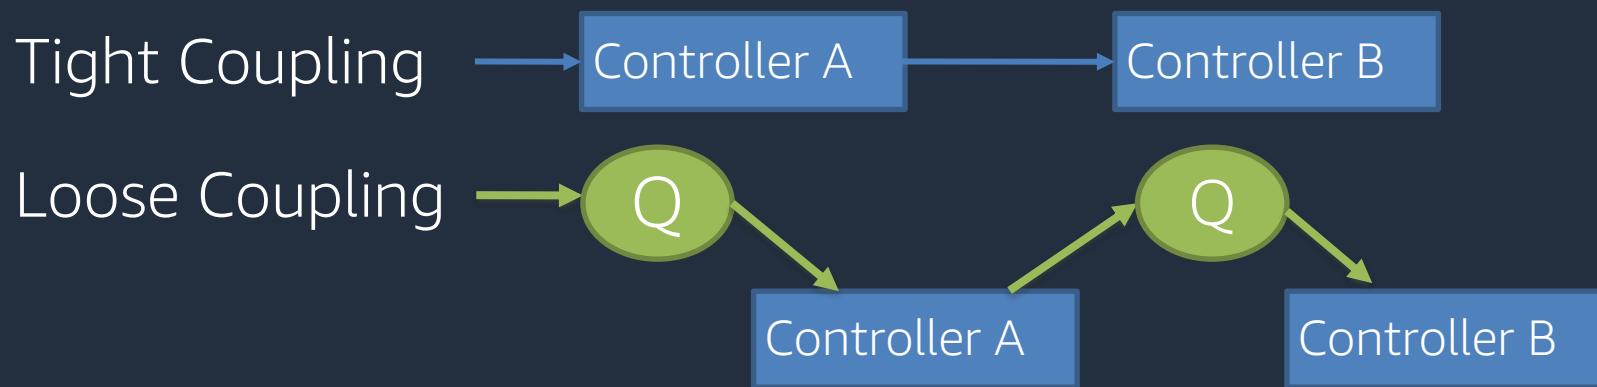
- Decompose Databases
- **Database per microservice pattern**

Allows easy use of **Canary** and **Blue-Green** deployments

# Key Elements of Microservices...

Each microservice is:

- **Elastic**: scales up or down independently of other services
- **Resilient**: services provide fault isolation boundaries
- **Composable**: uniform APIs for each service
- **Minimal**: highly cohesive set of entities
- **Complete**: loosely coupled with other services



# Microservices Benefits

Fast to develop

Rapid deployment

Parallel development & deployment

Closely integrated with DevOps

- Now "DevSecOps"

Improved scalability, availability & fault tolerance

More closely aligned to business domain

# Principles of the Two Pizza Team



- Two-pizza teams
- Full ownership
- Full accountability
- Aligned incentives
- “DevOps”

# How do Two Pizza Teams work?

We call them “Service teams”

Owning the “primitives” they build:

- Product planning (roadmap)
- Development work
- Operational/Client support work

**“You build it, you run it”**

Part of a larger concentrated org (Amazon.com, AWS, Prime, etc)

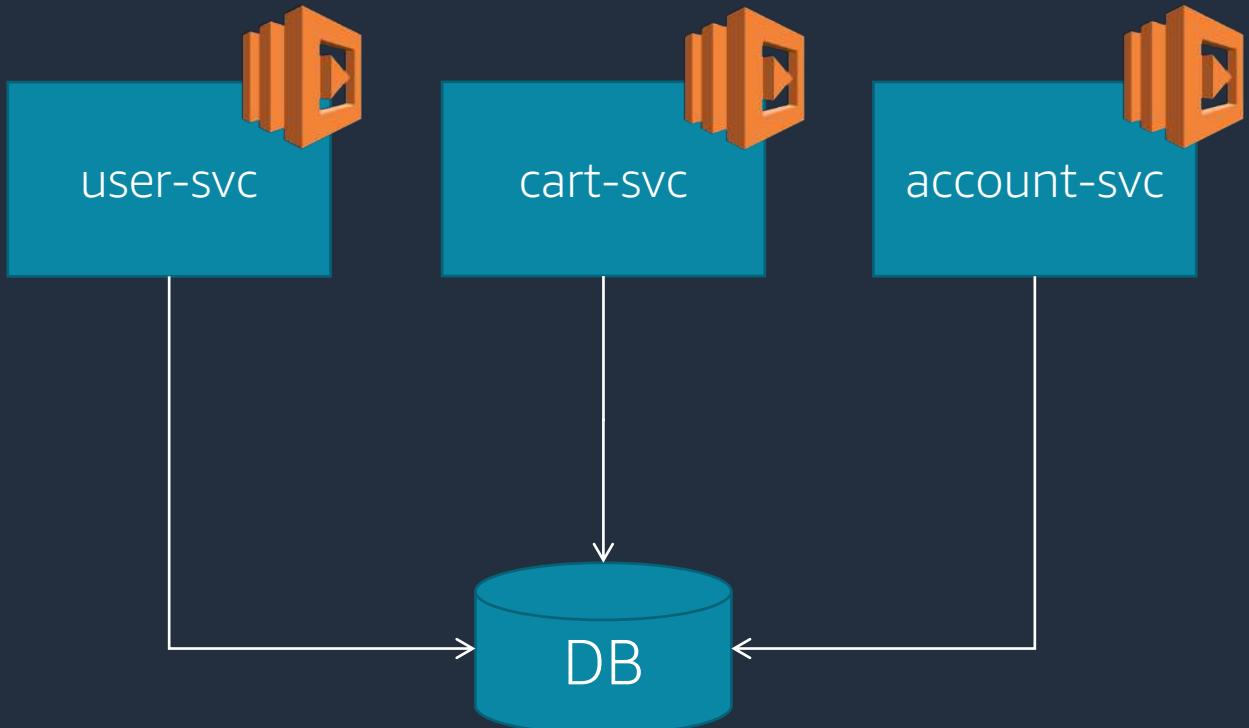


# Data Architecture Challenges

# Challenge: Centralized Database

Applications often have a **monolithic** data store

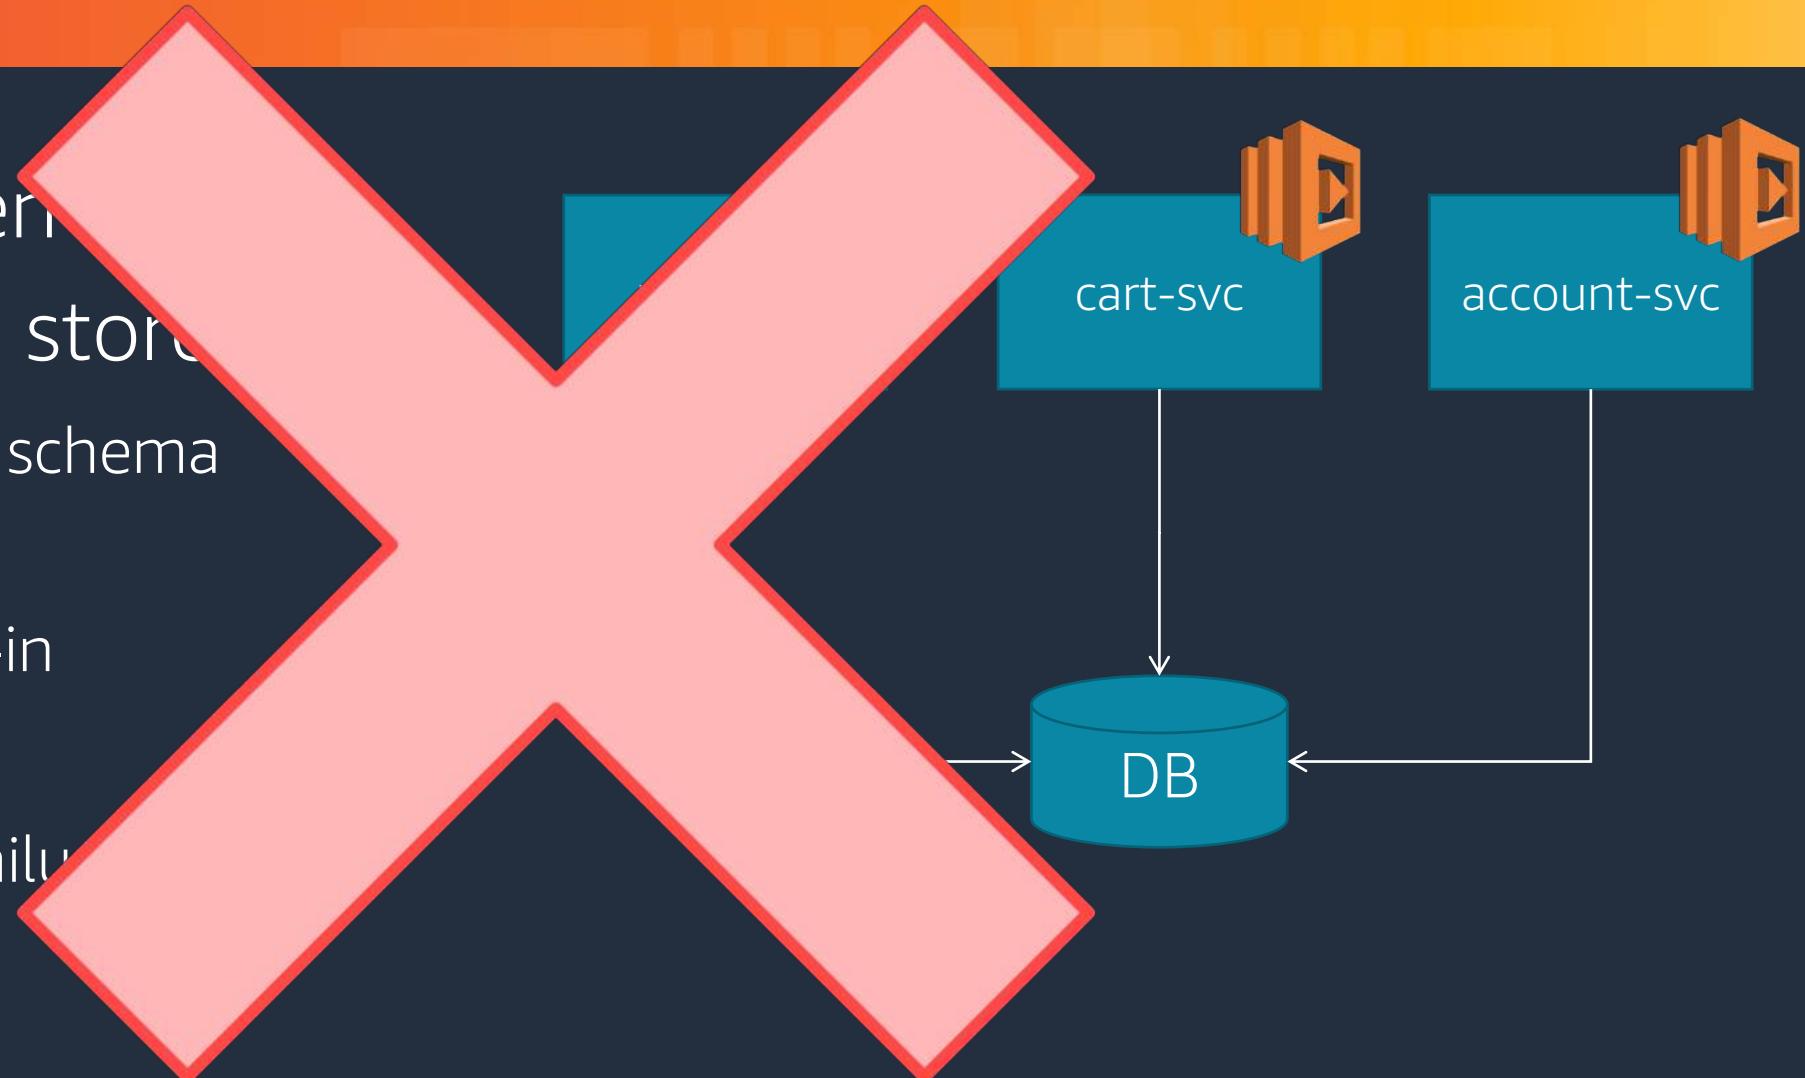
- Difficult to make schema changes
- Technology lock-in
- Vertical scaling
- Single point of failure



# Centralized Database – Anti-pattern

Applications often use a single, central database to store all data, known as a **monolithic** data store.

- Difficult to make schema changes
- Technology lock-in
- Vertical scaling
- Single point of failure



# Decentralized Data Stores

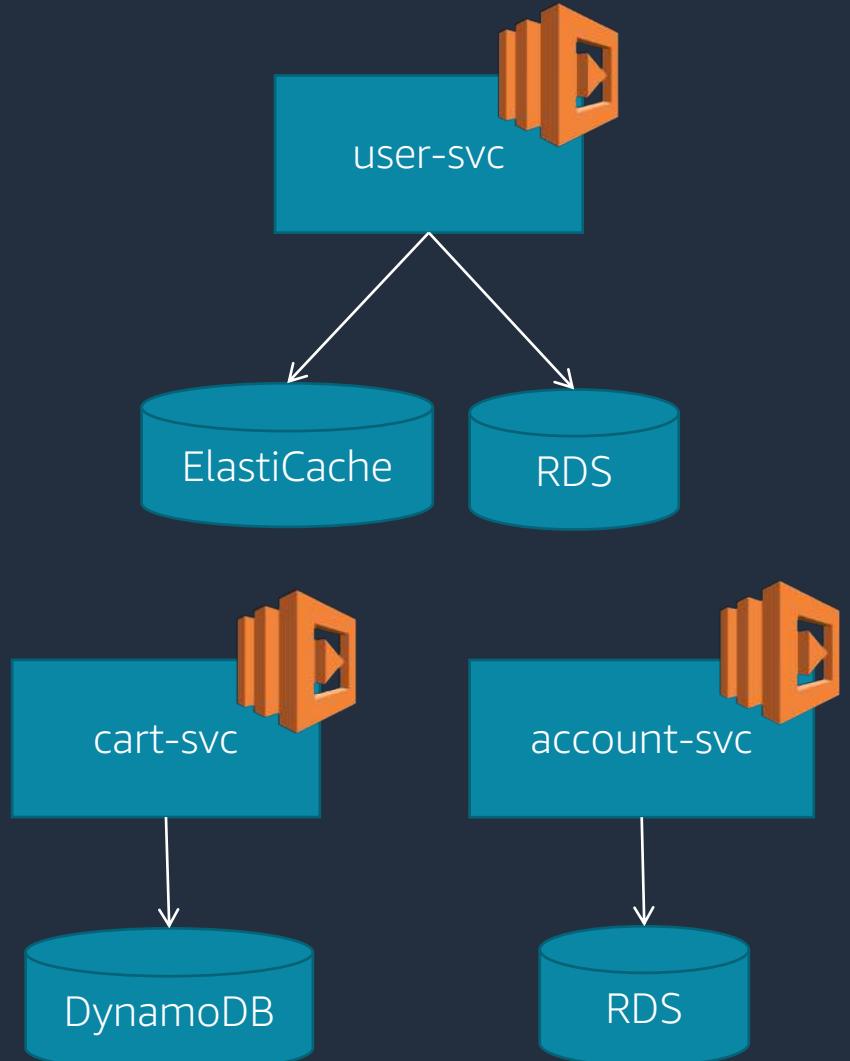
## Polyglot Persistence

Each service chooses it's data store technology

Low impact schema changes

Independent scalability

Data is gated **through the service API**



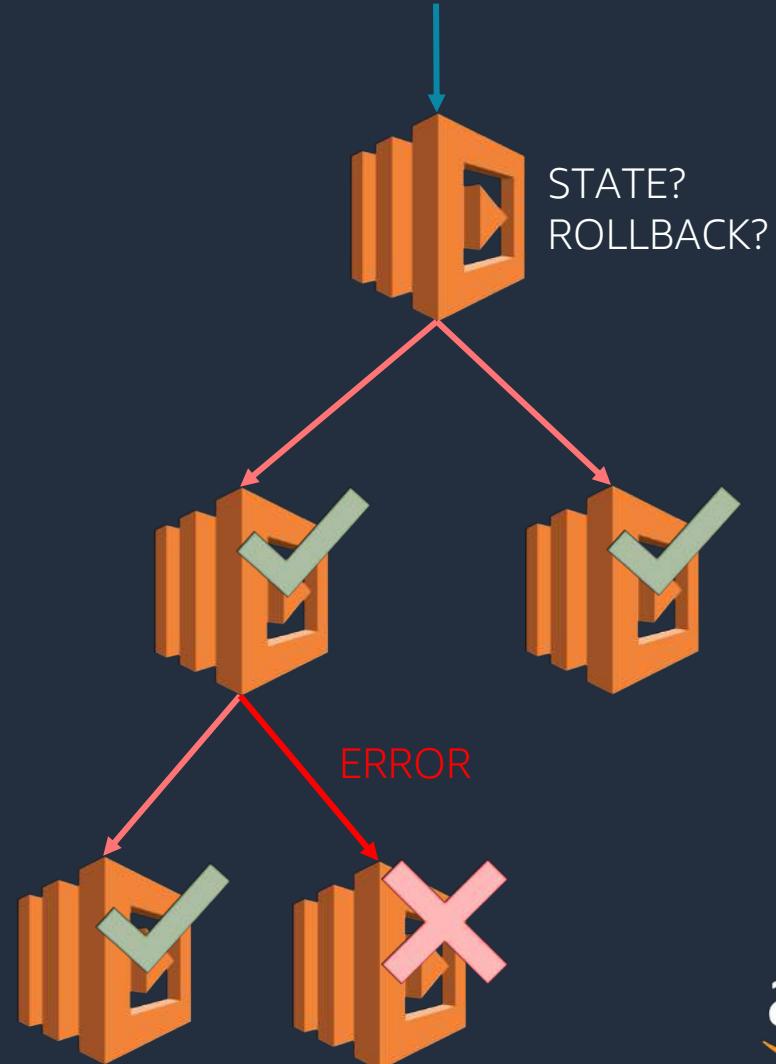
# Challenge: Transactional Integrity

Polyglot persistence generally translates into  
**eventual consistency**

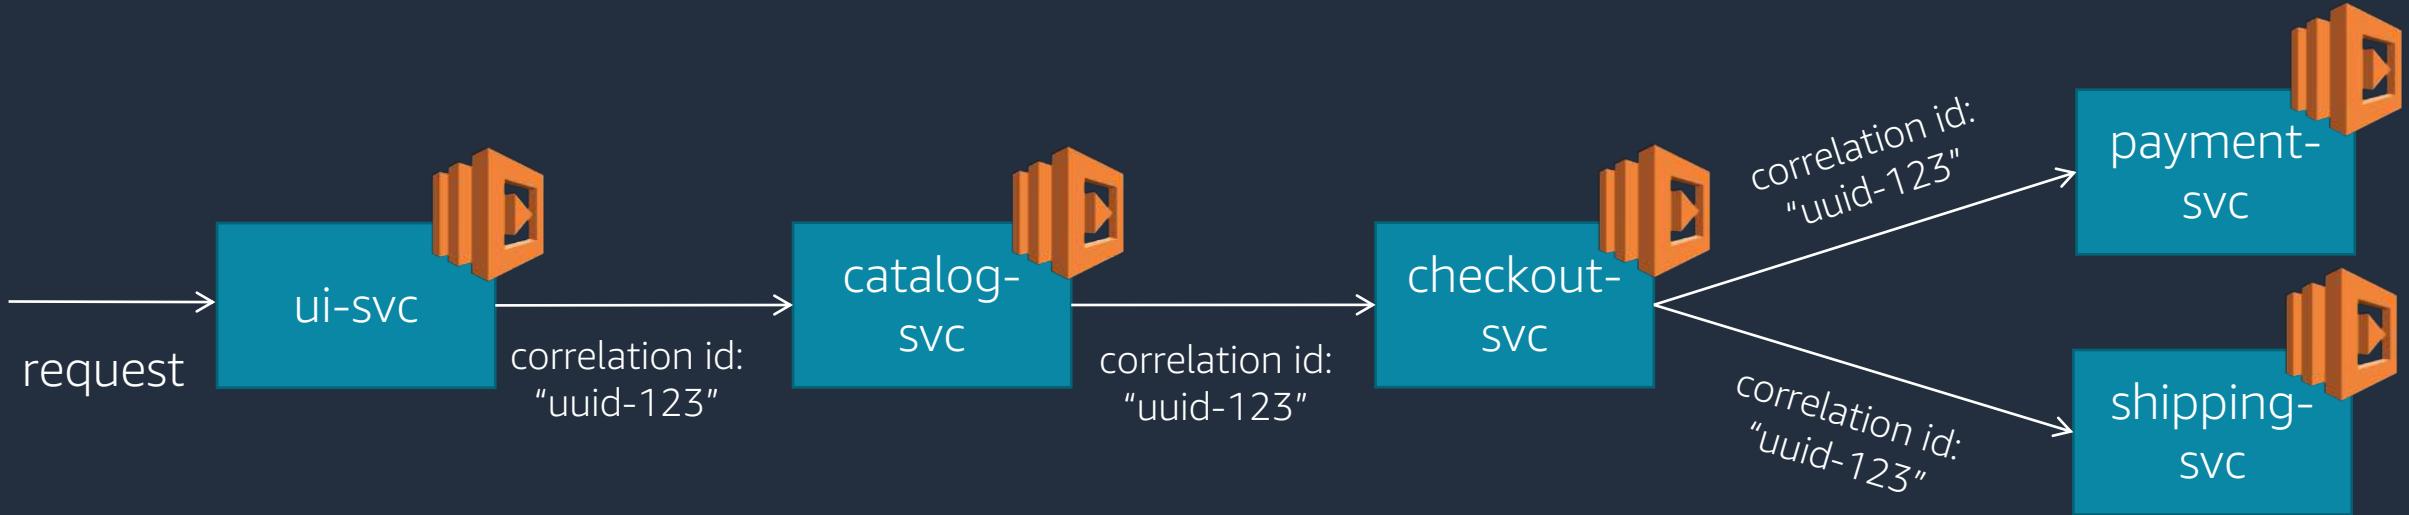
**Asynchronous calls** allow non-blocking, but  
returns need to be handled properly

How about **transactional integrity**?

- Event-sourcing – Capture changes as sequence of events
- Staged commit
- Rollback on failure



# Best Practice: Use Correlation IDs



```
09-02-2015 15:03:24 ui-svc INFO [uuid-123] ....  
09-02-2015 15:03:25 catalog-svc INFO [uuid-123] ....  
09-02-2015 15:03:26 checkout-svc ERROR [uuid-123] ....  
09-02-2015 15:03:27 payment-svc INFO [uuid-123] ....  
09-02-2015 15:03:27 shipping-svc INFO [uuid-123] ....
```



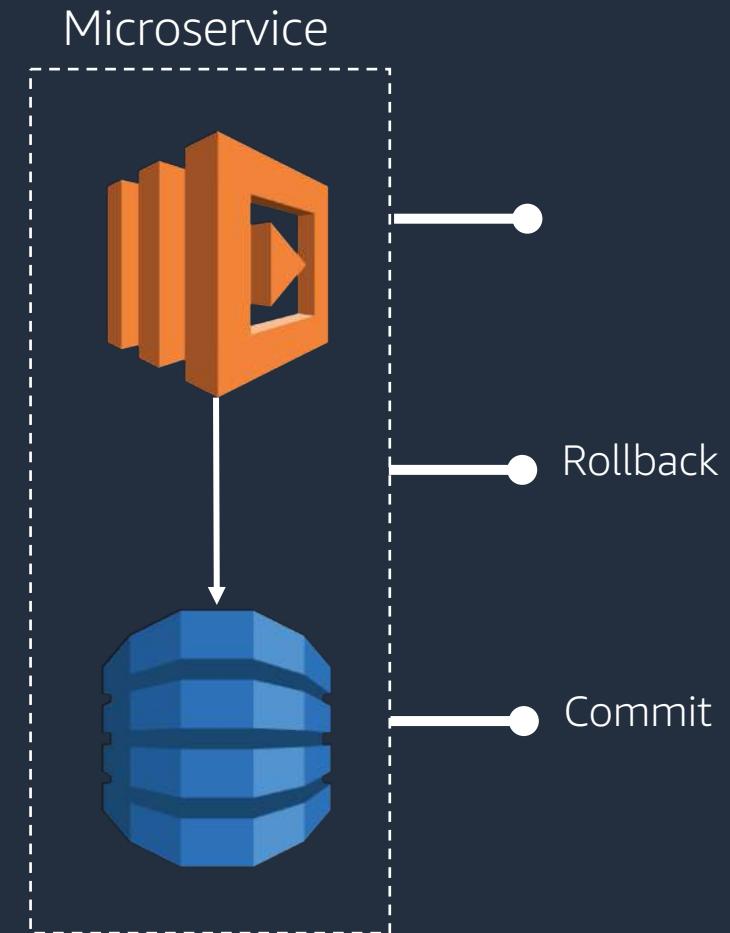
# Best Practice: Microservice owns Rollback

Every microservice should expose it's own "rollback" method

This method could just **rollback** changes, or trigger **subsequent actions**

- Could send a notification

If you implement **staged commit**, also expose a commit function



# Event-Driven: DynamoDB Streams

If async, consider event-driven approach with **DynamoDB Streams**

Don't need to manage function execution failure, DDB Streams

**automatically retries** until successful

**“Attach”** yourself to the data of interest



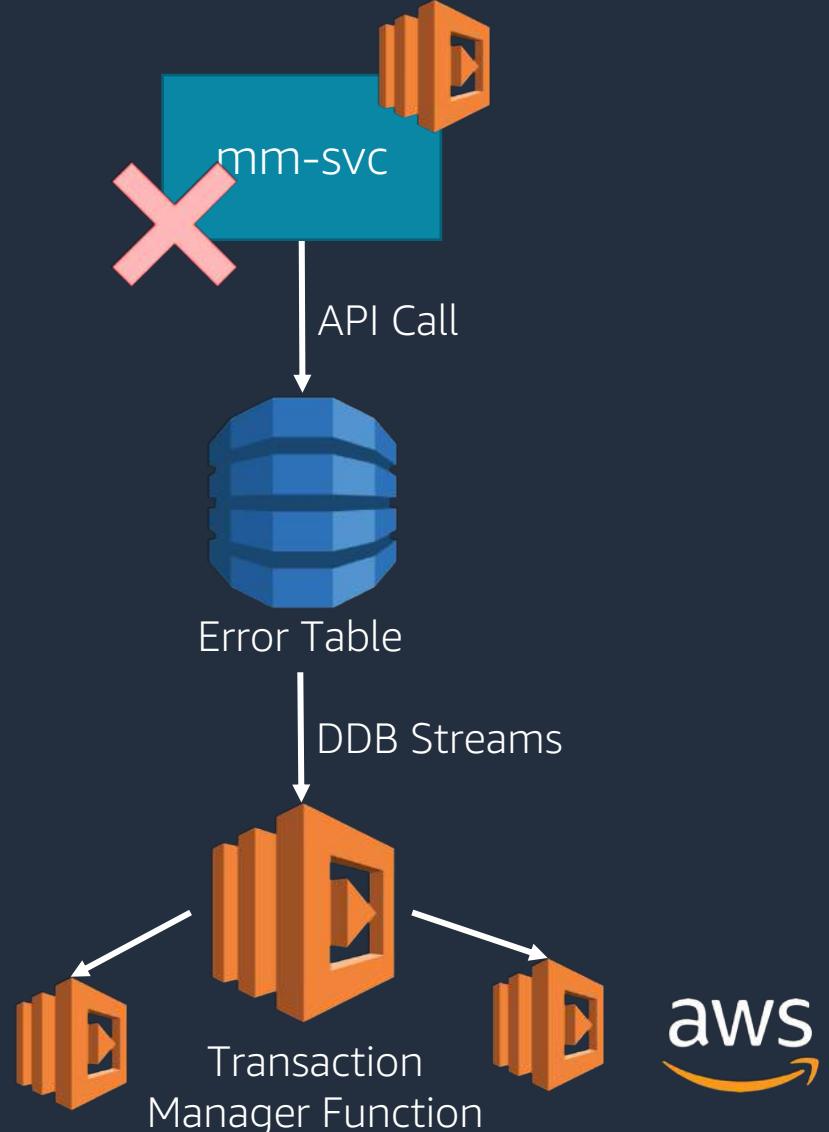
# Challenge: Report Errors / Rollback

What if functions fail? (business logic failure, not code failure)

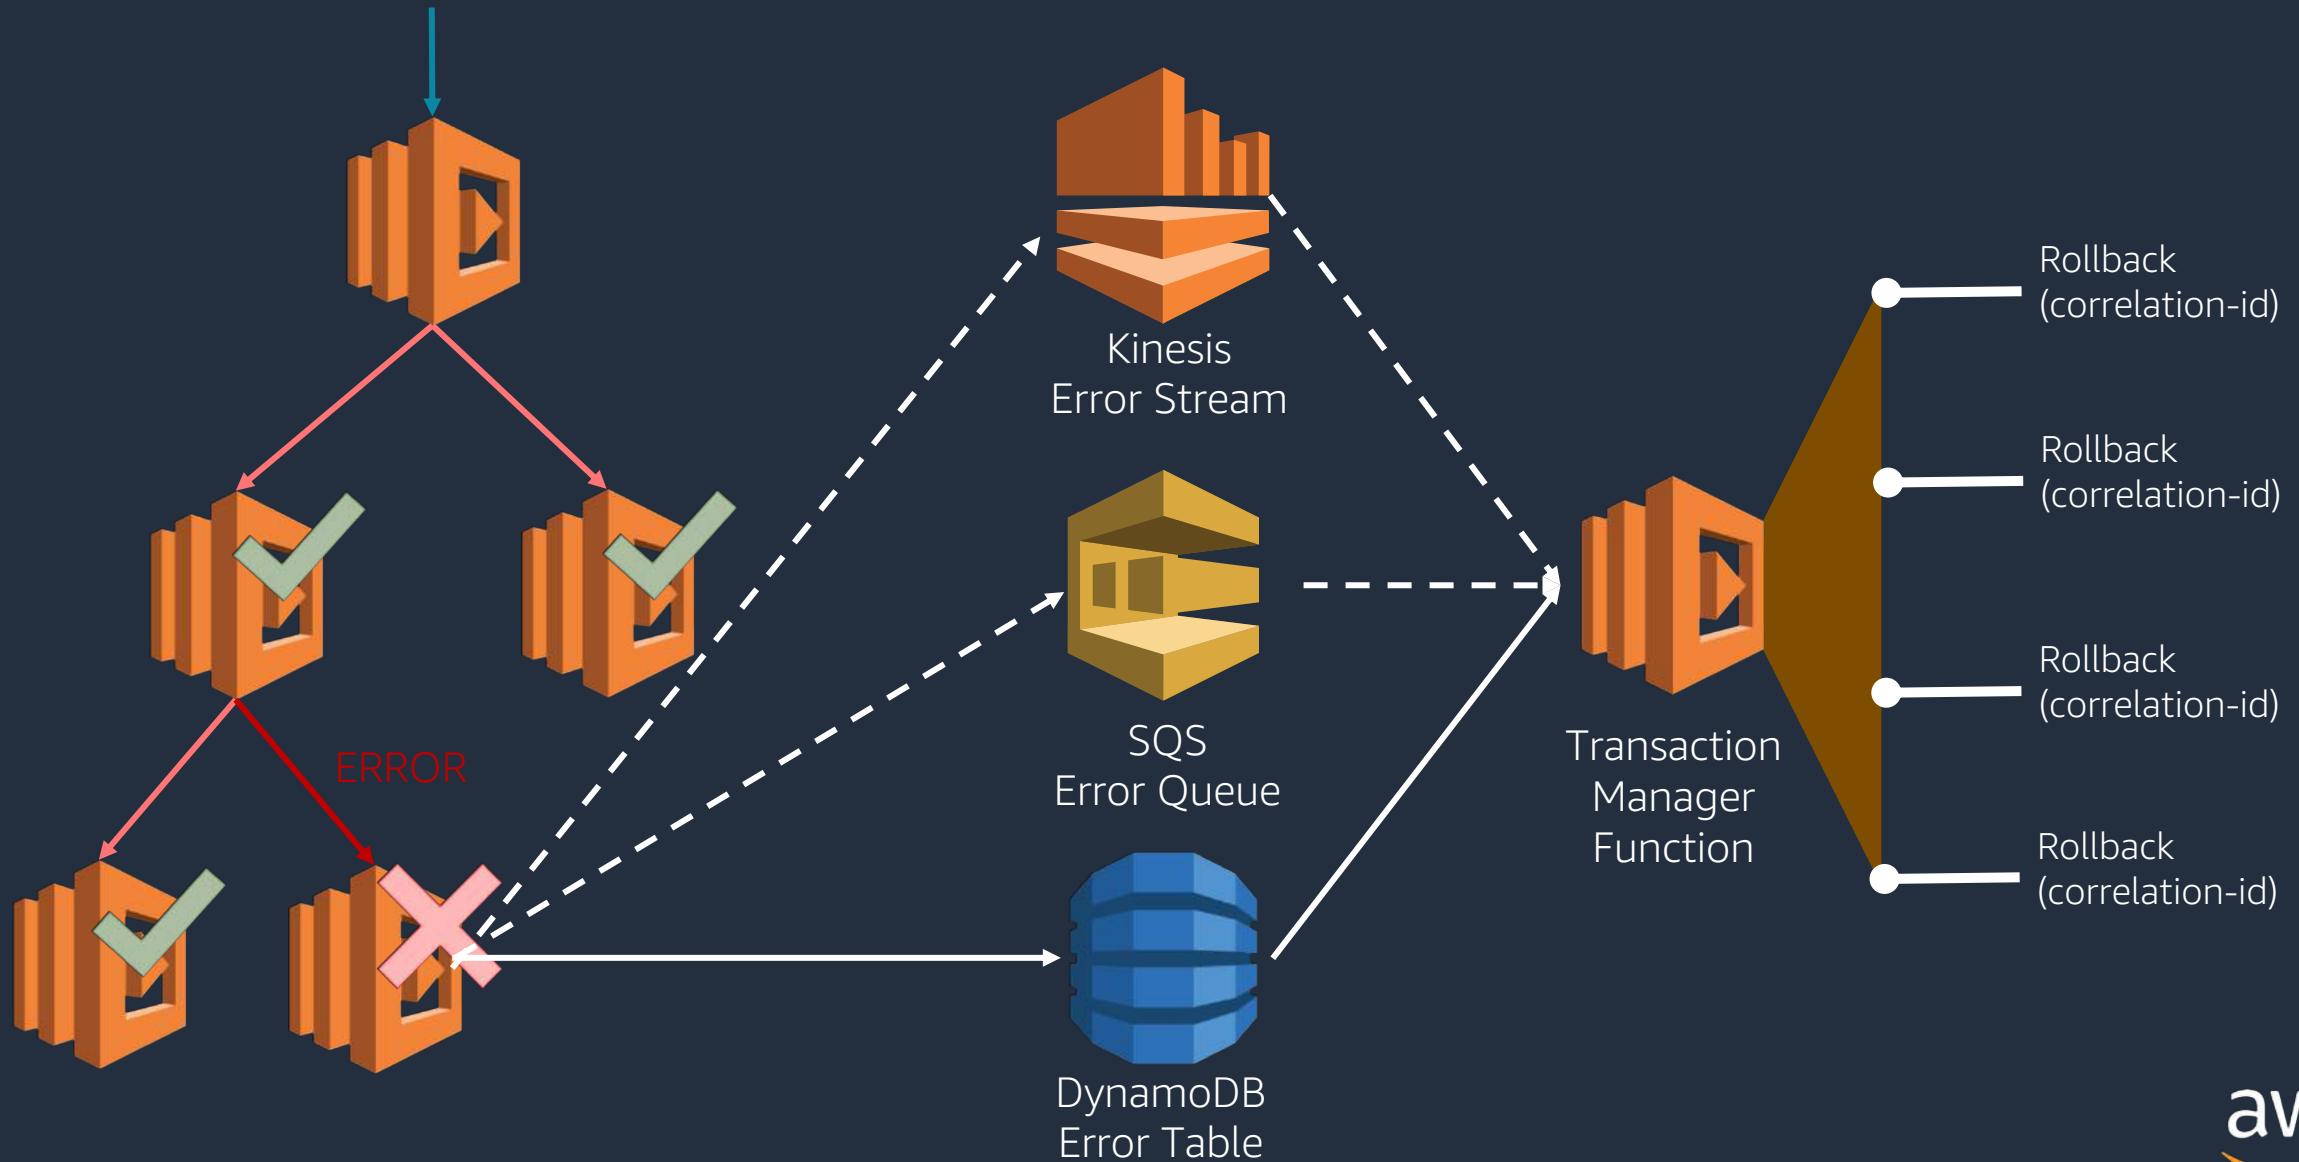
Create a **“Transaction Manager”** microservice that notifies all relevant microservices to rollback or take action

DynamoDB is the **trigger** for the clean-up function (could be SQS, Kinesis etc.)

Use **Correlation ID** to identify relations



# Challenge: Report Errors / Rollback

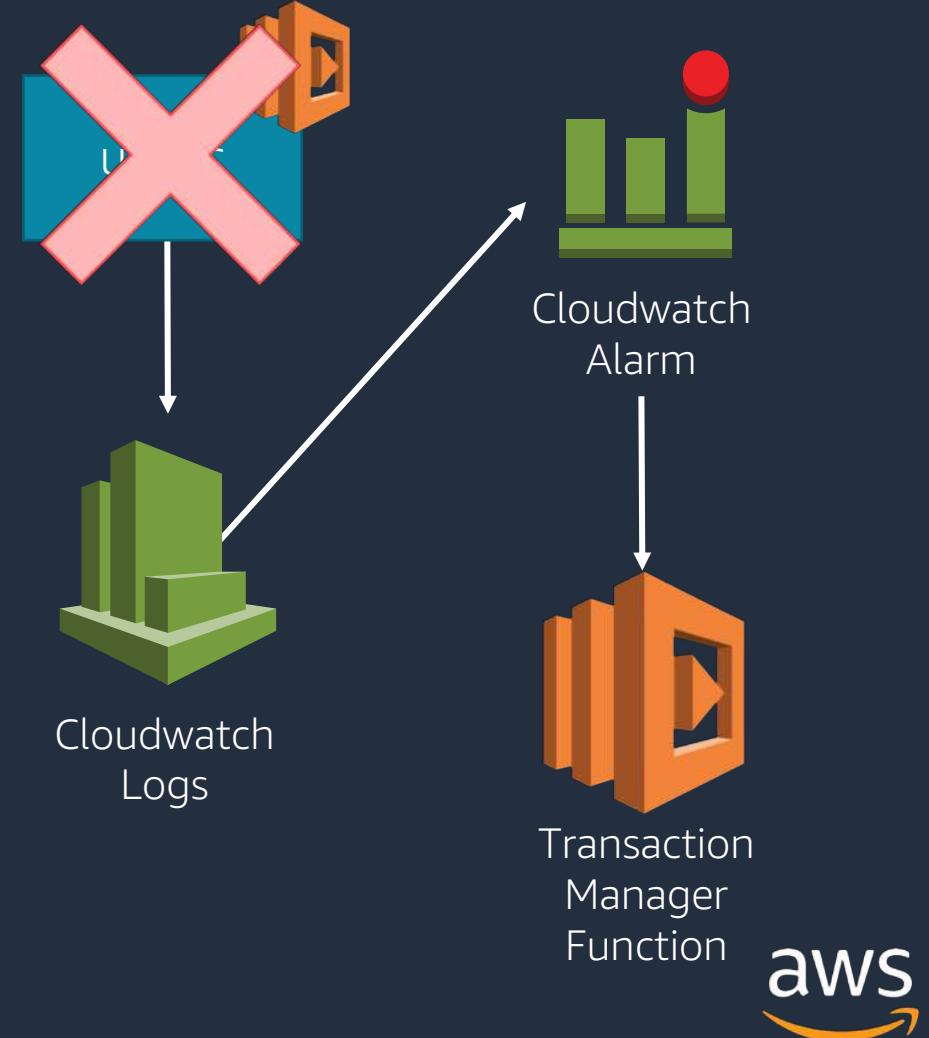


# Challenge: Code Error

Lambda Execution Error because of  
**faulty code**

Leverage **Cloudwatch Logs** to process  
error message and call Transaction  
Manager

Set **Cloudwatch Logs Metric Filter** to  
look for Error/Exception and call Lambda  
Handler upon Alarm state

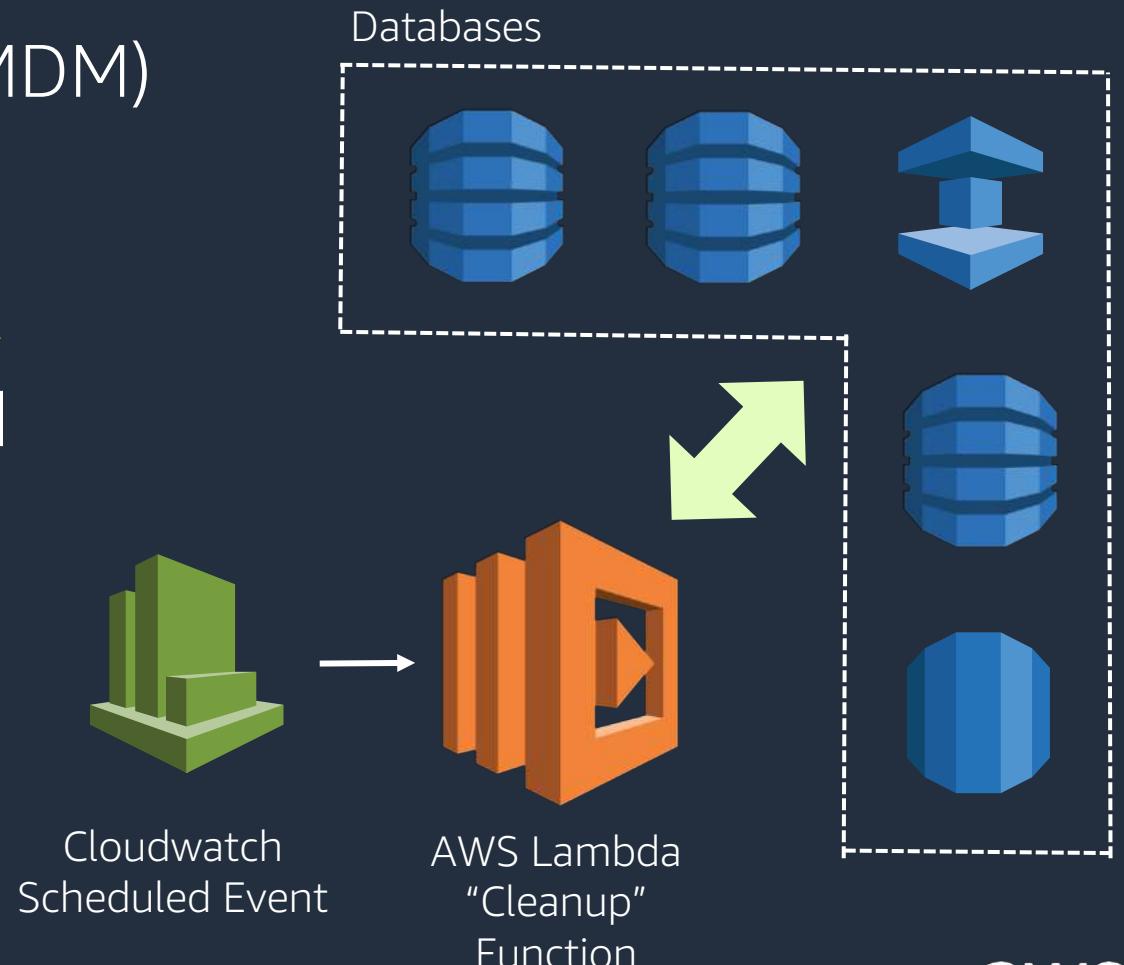


# MDM – Keep Data Consistent

Perform **Master Data Management** (MDM) to keep data consistent

Create AWS Lambda function to **check consistencies** across microservices and "cleanup"

Create a **Cloudwatch Event** to schedule the function (e.g. hourly basis)



# Choosing a Datastore

# Storage & DB options in AWS

In-Memory      NoSQL      SQL      Graph      Object      Search      Streaming



Amazon  
ElastiCache



Amazon  
DynamoDB



Amazon  
RDS



Amazon  
Redshift



Amazon  
Neptune



Amazon  
S3



Amazon  
Glacier

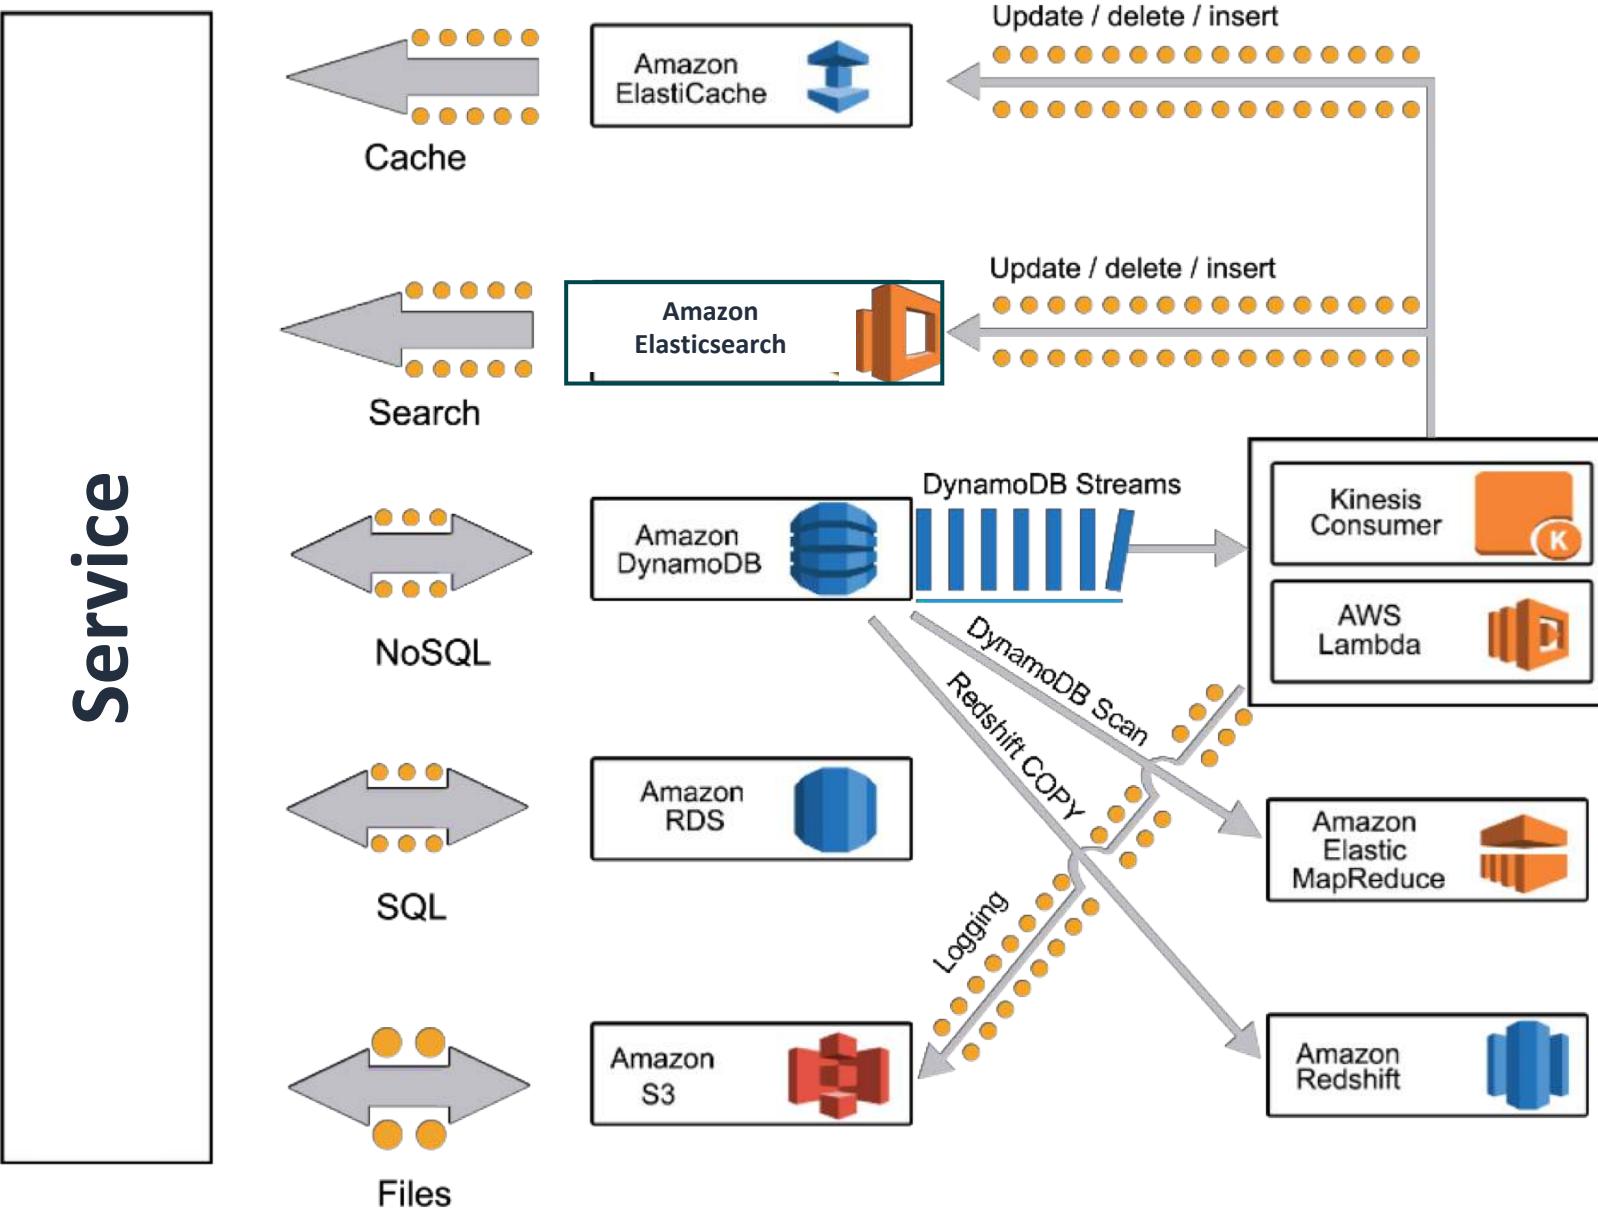


Amazon  
Elasticsearch  
Service



Amazon  
Kinesis





# Challenge: What Service to Use?

- Many problems can be solved with NoSQL, RDBMS or even in-memory cache technologies
- Non-functional requirements can help identify appropriate services
- **Solution:** Classify your organizations non-functional requirements and map them to service capabilities

# Determine Your Non-Functional Requirements

This is only an example. Your company's classifications will be different

Requirement				
<b>Latency</b>	> 1s	200 ms - 1s	20 ms – 200 ms	< 20 ms
<b>Durability</b>	99.99	99.999	99.9999	> 99.9999
<b>Storage Scale</b>	< 256 GB	256 GB – 1 TB	1 TB – 16 TB	> 16 TB
<b>Availability</b>	99	99.9	99.95	> 99.95
<b>Data Class</b>	Public	Important	Secret	Top Secret
<b>Recoverability</b>	12 – 24 hours	1 – 12 hours	5 mins – 1 hour	< 5 mins
<b>Skills</b>	None	Average	Good	Expert

There will be other requirements such as regulatory compliance.



# Map Non-Functional Requirements to Services

The information below is not exact and does not represent SLAs

Service	Latency	Durability	Storage	Availability	Recoverability from AZ Failure (RPO, RTO)
<b>RDS</b>	< 100 ms	> 99.8 (EBS)	16 TB	99.95	0s and 90s (MAZ)
<b>Aurora</b>	< 100 ms	> 99.9	64 TB	> 99.95	0s and < 30s (MAZ)
<b>Aurora + ElastiCache</b>	< 1 ms	> 99.9	64 TB	> 99.95	0s and < 30s (MAZ)
<b>DynamoDB</b>	< 10 ms	> 99.9	No Limit	> 99.99	0s and 0s
<b>DynamoDB / DAX</b>	< 1 ms	> 99.9	No Limit	> 99.99	0s and 0s
<b>ElastiCache Redis</b>	< 1 ms	N/A	3.5 TiB	99.95	0s and < 30s (MAZ)
<b>Elasticsearch</b>	< 200 ms	> 99.9	150 TB	99.95	0s and < 30s (Zone Aware)
<b>S3</b>	< 500 ms	99.999999999	No Limit	99.99	0s and 0s

# Finalizing Your Data Store Choices

- After mapping your non-functional requirements to services you should have a short list to choose from
- Functional requirements such as geospatial data and query support will refine the list further
- You may institute standards to make data store selection simpler and also make it easier for people to move between teams, e.g Redis over Memcached and PostgreSQL over MySQL. These can still be overridden, but require justification to senior management

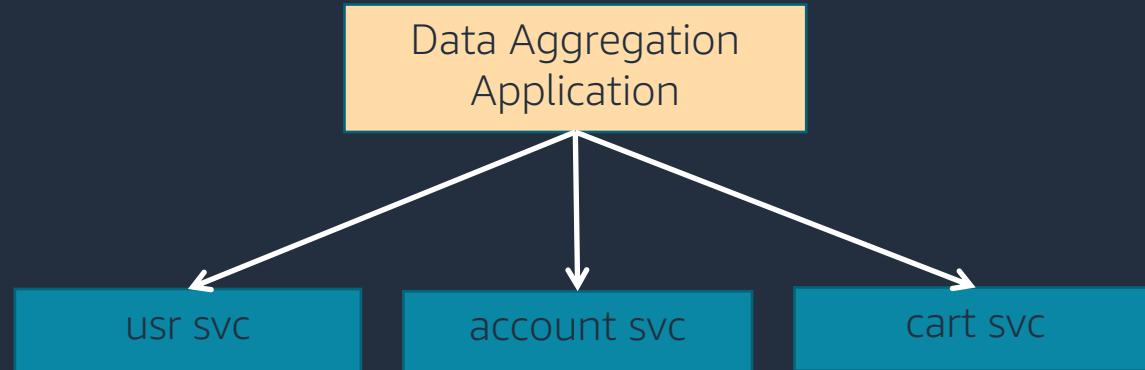


# Challenge: Reporting and Analytics

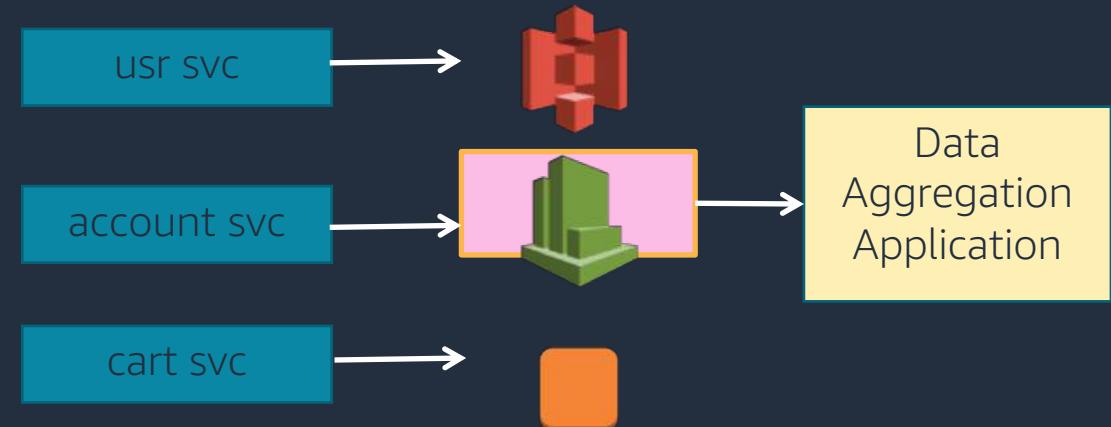
- Data is now spread across a number of isolated polyglot data stores
- Consolidation and aggregation required
- **Solution:** Pull data from required microservices, push data to data aggregation service, use pub/sub. Don't use a composite service (anti-pattern).

# Aggregation

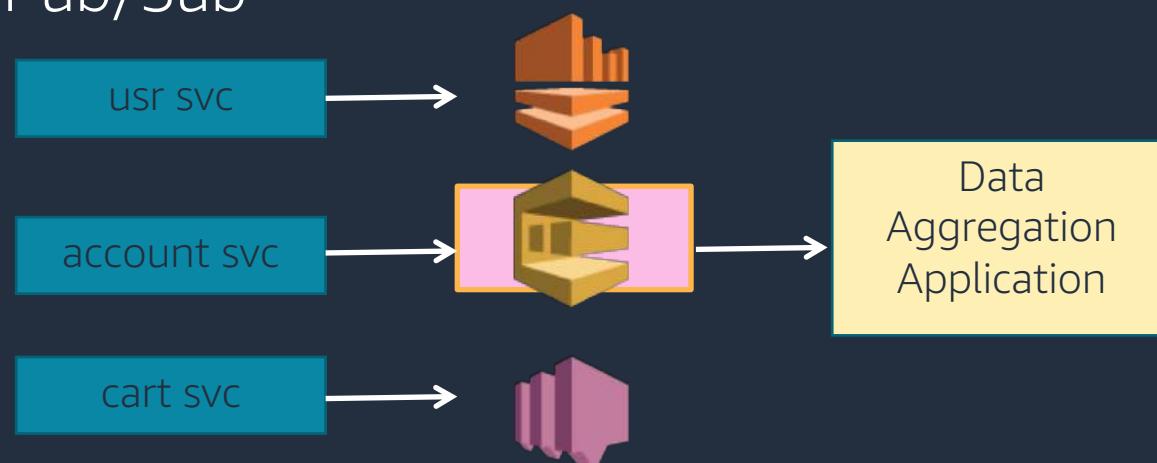
## Pull model



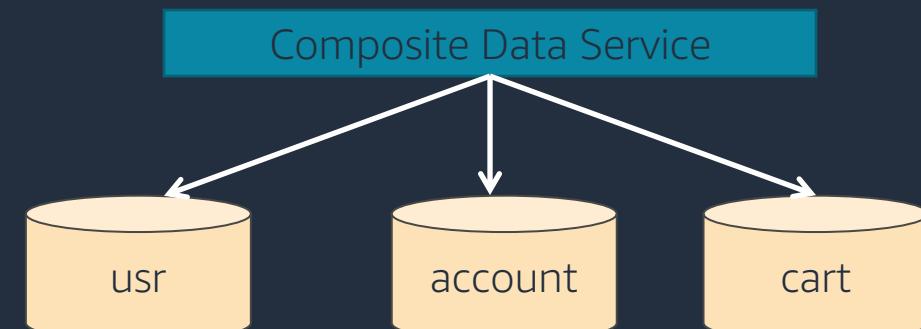
## Push model



## Pub/Sub



## Composite



# A Few Thoughts

- Use **Non-Functional Requirements** to help identify the right data store(s) for each microservice
- Use **polyglot persistence** to avoid bottlenecks, schema issues and allow independent scalability (and cache)
- Embrace **eventual consistency** and design fault-tolerant business processes which can recover
- Think ahead and plan your **analytics requirements** as part of the overall architecture



# INNOVATE2018

ONLINE CONFERENCE

DEVELOPER EDITION

# Thank you!

@awscloud



#AWSInnovate