



INNOVATE2018
ONLINE CONFERENCE

DEVELOPER EDITION

Serverless Architectural Patterns

Julio Faerman @faermanj

@awscloud



#AWSInnovate



NDC { Oslo }

5

Motivation or Hype? Time & Money



Log in to your 2016 Census



Thank you for participating in the Census. The system is very busy at the moment. Please wait for 15 minutes before trying again. Your patience and cooperation are appreciated. [code 9]

\$9 Million vs. \$ 500

A spectrum of managed services

"On EC2"



Amazon EC2



Microsoft SQL
Server



kafka



docker



cassandra

Managed



Amazon
EMR



Amazon ES



Amazon
ElastiCache



Amazon
Redshift



Amazon
RDS

Serverless



AWS
Lambda



Amazon
Cognito



Amazon
Kinesis



Amazon
S3



Amazon
DynamoDB



Amazon
SQS



Amazon API
Gateway



Amazon
CloudWatch



AWS IoT



Let's talk about...

Event Processing Architecture
Operations Automation Architecture
Web Application Architecture
Stream Processing Architecture
Telemetry Processing Architecture



Serverless means...



**No servers to provision
or manage**



Scales with usage



Never pay for idle

**Availability and fault
tolerance built in**



Regional services

Serverless Service

AZ1

AZ2

AZ3



Anatomy of a Lambda function

Handler() function

Function to be executed upon invocation

Event object

Data sent during Lambda Function Invocation

Context object

Methods available to interact with runtime information (request ID, log group, etc.)

```
def handler(event, context):  
    return {  
        "message": "Hello World!",  
        "event": event  
    }
```



Lambda execution model

Synchronous (push)



/api/hello



AWS Lambda
function

Asynchronous (event)



reqs



AWS Lambda
function

Stream-based



changes



AWS Lambda
service



function




Lambda Best Practices


- Minimize package size to necessities
- Separate the Lambda handler from core logic
- Use Environment Variables to modify operational behavior
- Self-contain dependencies in your function package
- Leverage “Max Memory Used” to right-size your functions
- Delete large unused functions (75GB limit)

AWS X-Ray Integration with Serverless


- Lambda instruments incoming requests for all supported languages
- Lambda runs the X-Ray daemon on all languages with an SDK



```
var AWSXRay = require('aws-xray-sdk-core');
AWSXRay.middleware.setSamplingRules('sampling-rules.json');
var AWS = AWSXRay.captureAWS(require('aws-sdk'));
S3Client = AWS.S3();
```



X-Ray Trace Example



Serverless Frameworks



Chalice



[awslabs/aws-serverless-express](#)

[awslabs/aws-serverless-java-container](#)



Event Processing Architecture

Event driven

Event A on B triggers C



Invocation



Action



Event-driven platform

Invoked in response to events

- Changes in data
- Changes in state



S3 event notifications



DynamoDB Streams



Kinesis events



SNS events



CloudTrail events



Cognito events



Custom events



CloudWatch events



Lambda functions

Access any service, including your own

Any custom



Such as...



SNS



DynamoDB



Lambda



Redshift




Kinesis




S3



Event-driven actions



Event-driven controls





AWS Step Functions:

Orchestrate a Serverless processing workflow using AWS Lambda

Manages state, checkpoints and restarts for you to make sure that your application executes in order and as expected

Build visual workflows that enable fast translation of business requirements





<https://github.com/awslabs/lambda-refarch-imagerecognition>

Operations Automation Architecture


Automation characteristics

- Periodic jobs
- Event triggered workflows
- Enforce security policies
- Audit and notification
- Respond to alarms
- Extend AWS functionality

... All while being Highly Available, Scalable and Auditable




Auto tagging resources as they start




Scheduled backup operation



Scheduled AI jobs



AWS Ops Automator




<https://aws.amazon.com/answers/infrastructure-management/ops-automator/>



Web Application Architecture




Web application



Serverless web app security



Custom Authorizers



Two types:

- TOKEN - authorization token passed in a header
- REQUEST – all headers, query strings, paths, stage variables or context variables.



Bustle Achieves 84% Cost Savings with AWS Lambda

"

With AWS Lambda, we
eliminate the need to worry
about operations

Tyler Love
CTO, Bustle



"

Bustle is a news, entertainment, lifestyle, and fashion website targeted towards women.

- Bustle had trouble scaling and maintaining high availability for its website without heavy management
- Moved to serverless architecture using AWS Lambda and Amazon API Gateway
- Experienced approximately 84% in cost savings
- Engineers are now focused on innovation





Stream Processing Architecture

Kinesis-based apps



Amazon Kinesis makes it easy to work with real-time streaming data



Amazon Kinesis Streams

- For Technical Developers
- Collect and stream data for ordered, replay-able, real-time processing

Amazon Kinesis Analytics


- For all developers, data scientists
- Easily analyze data streams using standard SQL queries

Amazon Kinesis Firehose

- For all developers, data scientists
- Easily load massive volumes of streaming data into Amazon S3, Redshift, ElasticSearch




Amazon Kinesis




** A *shard* is a group of data records in a stream

Amazon Kinesis under the hood



[OUR STORY](#)[CAREERS](#)[GAMES](#)[SUPPORT](#)[CONTACT US](#)[SUPERCELL SHOP](#)

CLASH of CLANS



Netflix & Amazon Kinesis Streams Case Study



Amazon Kinesis Streams processes multiple terabytes of log data each day, yet events show up in our analytics in seconds. We can discover and respond to issues in real time, ensuring high availability and a great customer experience.

John Bennett

Senior Software Engineer, Netflix

About Netflix

Netflix is the world's leading internet television network, with more than 100 million members in more than 190 countries enjoying 125 million hours of TV shows and movies each day.

AWS Services Used

Benefits of AWS

About Netflix


Netflix is the world's leading internet television network, with more than 100 million members worldwide enjoying 125 million hours of TV shows and movies each day, including original series, documentaries, and feature films. Members can watch as much as they want, anytime, anywhere, on nearly any Internet-connected screen.

Application Monitoring on a Massive Scale


Netflix uses Amazon Web Services (AWS) for nearly all its computing and storage needs, including databases, analytics, recommendation engines, video transcoding, and more—hundreds of functions that in total use more than 100,000 server instances on AWS.

This results in an extremely complex and dynamic networking environment where applications are constantly communicating inside AWS and across the Internet. Monitoring and optimizing its network is critical for Netflix to continue improving customer experience, increasing efficiency, and reducing costs. In particular, Netflix needed a solution for ingesting, augmenting, and analyzing the multiple terabytes of data its network generates daily in the form of virtual private cloud (VPC) flow logs. This would enable Netflix to identify performance-improvement opportunities, such as identifying apps that are communicating across regions and collocating them. The company would also be able to increase uptime by quickly detecting and mitigating application downtime.

Real-time analytics



Data Input



Source JSON Data

- Once per second, using JavaScript SDK:
 - Unique Cognito ID (anonymous user)
 - OS
 - Quadrant
 - Data sent to Kinesis Stream

```
{  
  "recordTime": 1486505943.204,  
  "cognitoId": "us-east-1:3626e211-d2a3-447b-8231-e1f4e0486f44",  
  "os": "Android",  
  "quadrant": "A",  
  ...  
}
```



How is raw data mapped to a schema?



Amazon Kinesis stream



Amazon Kinesis Analytics

```
{  
  "recordTime": 1486505943.204,  
  "cognitoId": "us-east-1:<guid>",  
  "os": "Android",  
  "quadrant": "A"  
}
```

cognitoID	os	quadrant
<guid1>	Android	A
<guid2>	iOS	B

Source Data for Kinesis Analytics

How is streaming data accessed with SQL?

STREAM

- Analogous to a TABLE
- Represents continuous data flow

```
CREATE OR REPLACE STREAM DISTINCT_USER_STREAM(  
    COGNITO_ID VARCHAR(64),  
    DEVICE VARCHAR(32),  
    OS VARCHAR(32),  
    QUADRANT char(1),  
    DT TIMESTAMP);
```




How is streaming data accessed with SQL?

PUMP

- Continuous INSERT query
- Inserts data from one in-application stream to another

```
CREATE OR REPLACE PUMP "DISTINCT_USER_PUMP" AS
  INSERT INTO "DISTINCT_USER_STREAM"
    SELECT STREAM DISTINCT
      "cognitoId",
      . . .
```

How do we get distinct user records?



Use PUMP to insert distinct records into in-app STREAM

```
CREATE OR REPLACE PUMP "DISTINCT_USER_PUMP" AS
  INSERT INTO "DISTINCT_USER_STREAM"
    SELECT STREAM DISTINCT
      "cognitoId",
      "device",
      "os",
      "quadrant",
      FLOOR(s.ROWTIME TO SECOND)
    FROM "SOURCE_SQL_STREAM_001" s;
```



How do we aggregate per second?


- Tumbling window, group by time period

```
CREATE OR REPLACE PUMP "OUTPUT_PUMP" AS
  INSERT INTO "DESTINATION_SQL_STREAM"
    SELECT STREAM
      COUNT(dus.COGNITO_ID) AS UNIQUE_USER_COUNT,
      COUNT((CASE WHEN dus.OS = 'Android' THEN COGNITO_ID ELSE null END)) AS ANDROID_COUNT,
      COUNT((CASE WHEN dus.OS = 'iOS' THEN COGNITO_ID ELSE null END)) AS IOS_COUNT,
      COUNT((CASE WHEN dus.OS = 'Windows Phone' THEN COGNITO_ID ELSE null END)) AS WINDOWS_PHONE_COUNT,
      COUNT((CASE WHEN dus.OS = 'other' THEN COGNITO_ID ELSE null END)) AS OTHER_OS_COUNT,
      COUNT((CASE WHEN dus.QUADRANT = 'A' THEN COGNITO_ID ELSE null END)) AS QUADRANT_A_COUNT,
      COUNT((CASE WHEN dus.QUADRANT = 'B' THEN COGNITO_ID ELSE null END)) AS QUADRANT_B_COUNT,
      COUNT((CASE WHEN dus.QUADRANT = 'C' THEN COGNITO_ID ELSE null END)) AS QUADRANT_C_COUNT,
      COUNT((CASE WHEN dus.QUADRANT = 'D' THEN COGNITO_ID ELSE null END)) AS QUADRANT_D_COUNT,
      ROWTIME
    FROM "DISTINCT_USER_STREAM" dus
   GROUP BY
     FLOOR(dus.ROWTIME TO SECOND);
```




Comparing Types of Windows

- Output created at the end of the window
- The output of the window will be single event based on the aggregate function used




Tumbling window
Aggregate per time interval



Sliding window
Windows constantly re-evaluated



Output to Kinesis Stream



```
{  
    "unique_user_count": 96,  
    "android_count": 50,  
    "ios_count": 46,  
    "android_count": 50,  
    "quadrant_a_count": 80,  
    "quadrant_b_count": 10,  
    "quadrant_c_count": 3,  
    "quadrant_d_count": 3  
}
```

Persist aggregated data in DynamoDB




```
event.Records.forEach((record) => {
  const payload = new Buffer(record.kinesis.data, 'base64').toString('ascii');
  var docClient = new AWS.DynamoDB.DocumentClient();
  var table = "user-quadrant-data";
  var data = JSON.parse(payload);
  var params = {
    TableName: table,
    Item: {
      "dataType": "quadrantRollup",
      "windowtime": (new Date(data.WINDOW_TIME)).getTime(),
      "userCount": data.UNIQUE_USER_COUNT,
      "quadrantA": data.QUADRANT_A_COUNT,
      "quadrantB": data.QUADRANT_B_COUNT, ...
    }
  };
  docClient.put(params, function(err, data) { ... }
```



Processing a Kinesis Streams with AWS Lambda

Kinesis Stream



Logs Analytics



<https://github.com/adhorn/logtoes>



DATABASE

logfree

TABLES

Add Table

Filter Tables...

logtofirehose

Demo LogToFirehose

```
1 CREATE EXTERNAL TABLE IF NOT EXISTS logfree.logtofirehose (
2   ip string,
3   datetime timestamp,
4   request string,
5   agent string,
6   user string,
7   status_code string,
8   raw_agent string
9 )
10 ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'
11 with serdeproperties( 'ignore.malformed.json' = 'true' )
12 LOCATION 's3://adhorn-logtofirehose/'
```

Run Query

Save As

Format Query

New Query

(Run time: 0.53 seconds, Data scanned: 0KB)

Amazon Athena

Results

Query successful.

Create Demo LogToFirehose

```
1 SELECT * FROM logfree."logtofirehose" order by datetime desc limit 100;
```

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Run Query

Save As

Format Query

New Query

(Run time: 1.11 seconds, Data scanned: 20.24KB)

...

Results



	ip	datetime	request	agent	user	status_code	raw_agent
1	127.0.0.1	2017-09-13 03:07:16.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
2	127.0.0.1	2017-09-13 03:07:14.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
3	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
4	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
5	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
6	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
7	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
8	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
9	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
10	127.0.0.1	2017-09-13 03:07:11.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
11	127.0.0.1	2017-09-13 03:07:13.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
12	127.0.0.1	2017-09-13 03:07:12.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
13	127.0.0.1	2017-09-13 03:07:12.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36
14	127.0.0.1	2017-09-13 03:06:45.000	GET /api/echo	macos chrome 60.0.3112.113	guest	200	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/60.0.3112.113 Safari/537.36

Amazon Athena

Telemetry Processing Architecture


IoT-based apps




Sensor data collection



Anomaly Detection Using AWS Lambda



Interactive Serverless applications



Further Reading

Optimizing Enterprise Economics with Serverless Architectures

<https://d0.awsstatic.com/whitepapers/optimizing-enterprise-economics-serverless-architectures.pdf>

Serverless Architectures with AWS Lambda

<https://d1.awsstatic.com/whitepapers/serverless-architectures-with-aws-lambda.pdf>

Serverless Applications Lens - AWS Well-Architected Framework

<https://d1.awsstatic.com/whitepapers/architecture/AWS-Serverless-Applications-Lens.pdf>

Streaming Data Solutions on AWS with Amazon Kinesis

<https://d1.awsstatic.com/whitepapers/whitepaper-streaming-data-solutions-on-aws-with-amazon-kinesis.pdf>

AWS Serverless Multi-Tier Architectures

https://d1.awsstatic.com/whitepapers/AWS_Serverless_Multi-Tier_Archiectures.pdf



Key Takeaways

Event Processing Architecture

Operations Automation Architecture

Web Application Architecture

Stream Processing Architecture

Telemetry Processing Architecture



INNOVATE2018

ONLINE CONFERENCE

DEVELOPER EDITION

Thank you!

@awscloud



#AWSInnovate